

Tutorial1-RNN Sequence classifier

April 9, 2024

1 Tutorial 1 - RNN Sequence classifier

In this notebook, we will predict the winner of a basketball game based on the scores observed in the first 3 quarters of the game. Each column represents the beginning of a minute during the game. (There are 12 minutes in each quarter. There are 3 quarters in the data, so we have 36 columns as input variables in chronological order.) The values captured in each column represent the score difference observed at that minute (home score minus away score)

The last column W represents whether the home team (1) or the away team (0) won the game. This is the target variable. **Our unit of analysis is a single game.**

```
[ ]: # Common imports
import numpy as np
import tensorflow as tf
from tensorflow import keras
import pandas as pd
import matplotlib.pyplot as plt

np.random.seed(1)
tf.random.set_seed(1)

epoch_num = 5 # numer of epochs to train each model
```

1.1 Load and split data

Basketball data consists of each row (observation) representing a single game. The first 36 columns represent the score difference at each minute of the game. The last column represents the winner of the game (1 = home team, 0 = away team). The data is stored in a csv file. We will load the data into a pandas dataframe.

```
[ ]: data = pd.read_csv("https://raw.githubusercontent.com/prof-tcsmith/data/master/
↳basketball.csv")
```

We see that we have 1230 observations (games) and 37 columns (36 minutes of the game + 1 target variable).

```
[ ]: data.shape
```

```
[ ]: (1230, 37)
```

```
[ ]: data.head(5)
```

```
[ ]:      M1  M2  M3  M4  M5  M6  M7  M8  M9  M10  ...  M28  M29  M30  M31  M32  M33  \
0   -2  -1   1   1  -1  -3  -1   0   3    6  ...   9   11   10    7    7    4
1    0   2   7   6  10   8   8   6   0    6  ...   7   11   11   14   14   15
2    0  -2   2   0   5   4   5   3   5    3  ...   9   11   13   13   12   17
3    0   2   0   3   4   3   5   4   3    1  ...  10    6    7    7    8    8
4    0  -2  -2   0   3  -2  -7  -5  -7   -4  ...  10   10   15   13   11   11

      M34  M35  M36  W
0         6    2    1  1
1        13   13   13  0
2        15   15   12  1
3         8    8    8  1
4        11   13   10  1
```

[5 rows x 37 columns]

Split each observation into a sequence of 36 observations (minutes) and a single observation (winner). The sequence of 36 observations will be the input (X) to the RNN. The single observation (y) will be the target variable.

```
[ ]: y = data['W']
     X = data.drop('W', axis=1)
```

```
[ ]: y = np.array(y).astype(np.int8)
     X = np.array(X).astype(np.int8)
```

Finally, we need to split the data into training and test sets. A random split on this data is OK, as each 'game' (observation) is independent of the others. We will use 70% of the data for training and 30% for testing.

```
[ ]: from sklearn.model_selection import train_test_split

     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

1.2 Create dataframe to store results

We are running 8 models in this notebook. We will store the results in a dataframe for easy comparison.

```
[ ]: df_results = pd.DataFrame({'model': [], 'time': [], 'accuracy': [], 'precision':
    ↪ [], 'recall': [], 'f1': []})
```

1.3 MODEL01 - A normal (cross-sectional) NN

This model assumes that the data is NOT a time-series data set. It treats the data as cross-sectional and the columns being independent of each other.

```
[ ]: print(X_train.shape) # 2d array of shape (n_samples, n_features)
      print(y_train.shape) # 1d array of shape (n_samples,)
```

```
(861, 36)
(861,)
```

```
[ ]: model = keras.models.Sequential([
      keras.layers.Flatten(input_shape=[36, 1]),
      keras.layers.Dense(72, activation='relu'),
      keras.layers.Dense(1, activation='sigmoid')
    ])
```

```
/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
2024-04-09 07:30:26.758631: I metal_plugin/src/device/metal_device.cc:1154]
Metal device set to: Apple M1 Pro
2024-04-09 07:30:26.758655: I metal_plugin/src/device/metal_device.cc:296]
systemMemory: 16.00 GB
2024-04-09 07:30:26.758662: I metal_plugin/src/device/metal_device.cc:313]
maxCacheSize: 5.33 GB
2024-04-09 07:30:26.758680: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2024-04-09 07:30:26.758694: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)
```

```
[ ]: model.compile(
      loss="binary_crossentropy", # If multiclass, use
      ↪ "sparse_categorical_crossentropy" as the loss function
      optimizer=tf.keras.optimizers.Nadam(learning_rate=0.001),
      metrics=['accuracy']
    )

    model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-------------------|--------------|---------|
| flatten (Flatten) | (None, 36) | 0 |

```
dense (Dense)                (None, 72)                2,664
dense_1 (Dense)              (None, 1)                   73
```

Total params: 2,737 (10.69 KB)

Trainable params: 2,737 (10.69 KB)

Non-trainable params: 0 (0.00 B)

```
[ ]: ### How to interpret the output shape?
```

- The first layer
 - None is a wild-card representing the number of observations - this will be the number of observations in a batch that is fed into the network.
 - 36 is the number of units in the layer, since there will be 36 outputs from this layer 36, the second dimension will have 36
- The second layer
 - None is a wild-card representing the number of observations - this will be the number of observations in a batch that is fed into the network. This will be the same at each later, since the batch size is the same during each epoch.
 - 72 is the number of units in the layer, since there will be 72 outputs from these 72 units
- The third layer
 - None is a wild-card, is the number of observations given in the batch.
 - 1 is the number of units in the layer. Since the output is a binary value, there is one unit for the output, and therefore 1 is the number of outputs.

We can write this formulaically as:

```
[ ]: p = 1  # for NN, one weight per unit input
      b = 1  # for NN, one bias per unit

      # First hidden layer
      i = 36 # 36 inputs from the previous layer (flatten layer)
      h = 72 # number of units in this layer

      p * (h*i + b*h)
```

```
[ ]: 2664
```

1.3.1 Calculate the number of paramaters for the model

- The first later doesn't have any parameters, because it's just flattening the input
 - 0 inputs to 36 unites, is 0 weights and zero biases

- The second layer has 36 units connecting to 72 units, so 36×72 weights. If each weight also has a bias, then we add 72
 - $36 \times 72 + 72 = 2664$
- The third layer has 72 outputs from the previous layer to 1 unit, so 72 weights. If each unit in the layer has a bias, then we add 1 since there is only one unit
 - $72 + 1 = 73$

<https://towardsdatascience.com/counting-no-of-parameters-in-deep-learning-models-by-hand-8f1716241889>

```
[ ]: import time
start_time = time.time()

call_back = keras.callbacks.EarlyStopping(patience=5, restore_best_weights=True)
history = model.fit(
    X_train,
    y_train,
    epochs=epoch_num,
    validation_data=(X_test, y_test),
    callbacks=[call_back]
)

end_time = time.time()
```

Epoch 1/5

2024-04-09 07:30:27.617055: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:117]
Plugin optimizer for device_type GPU is enabled.

27/27 2s 25ms/step -
accuracy: 0.5047 - loss: 2.9070 - val_accuracy: 0.7371 - val_loss: 0.6694

Epoch 2/5

27/27 0s 11ms/step -
accuracy: 0.7198 - loss: 0.7530 - val_accuracy: 0.7805 - val_loss: 0.5617

Epoch 3/5

27/27 0s 11ms/step -
accuracy: 0.7477 - loss: 0.6374 - val_accuracy: 0.7805 - val_loss: 0.5219

Epoch 4/5

27/27 0s 11ms/step -
accuracy: 0.7589 - loss: 0.5625 - val_accuracy: 0.7751 - val_loss: 0.5007

Epoch 5/5

27/27 0s 12ms/step -
accuracy: 0.7667 - loss: 0.5233 - val_accuracy: 0.7940 - val_loss: 0.4889

```
[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))
fig.suptitle('NN Model Evaluation')

ax1.set_title('loss')
```

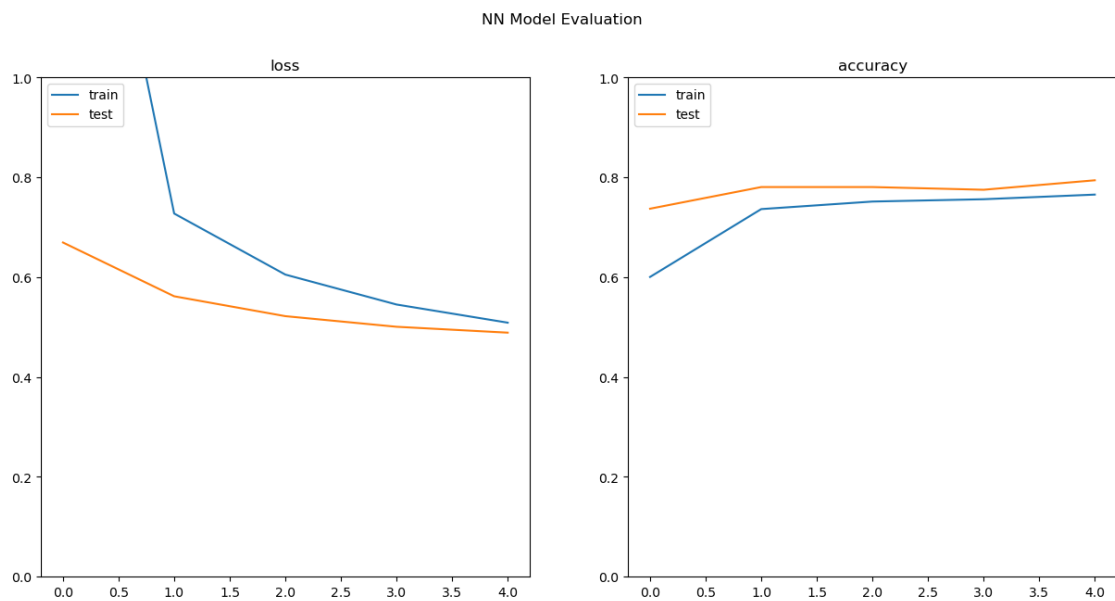
```

ax1.set_ylim(0, 1)
ax1.plot(history.history['loss'])
ax1.plot(history.history['val_loss'])
ax1.legend(['train', 'test'], loc='upper left')

ax2.set_title('accuracy')
ax2.set_ylim(0, 1)
ax2.plot(history.history['accuracy'])
ax2.plot(history.history['val_accuracy'])
ax2.legend(['train', 'test'], loc='upper left')

plt.show()

```



```

[ ]: y_pred = model.predict(X_test)
     y_pred[:10]

```

12/12 0s 5ms/step

```

[ ]: array([[1.2957262e-01],
           [8.8312477e-01],
           [2.1697335e-01],
           [6.0153572e-04],
           [2.4946399e-01],
           [5.0688112e-01],
           [3.8714585e-01],
           [9.5346004e-01],
           [2.1333496e-01],
           ...])

```

```
[9.5856142e-01]], dtype=float32)
```

```
[ ]: from sklearn.metrics import classification_report

y_pred = model.predict(X_test) # these predictions are probabilities (0-1)
    ↪ taken from the sigmoid function
y_pred = (y_pred > 0.5).astype('int8') # let's turn these into predictions (0
    ↪ or 1)

print(classification_report(y_test, y_pred, digits=4))
```

```
12/12          0s 1ms/step
      precision    recall  f1-score   support

     0       0.7407     0.8383     0.7865         167
     1       0.8500     0.7574     0.8010         202

 accuracy                   0.7940         369
  macro avg       0.7954     0.7979     0.7938         369
weighted avg       0.8006     0.7940     0.7945         369
```

Store the results into the results dataframe.

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↪ f1_score

df_result = pd.DataFrame(
    {
        'model': ['NN'],
        'time': [end_time - start_time],
        'accuracy': [accuracy_score(y_test, y_pred)],
        'precision': [precision_score(y_test, y_pred)],
        'recall': [recall_score(y_test, y_pred)],
        'f1': [f1_score(y_test, y_pred)]
    }
)

df_results = pd.concat([df_results, df_result], ignore_index=True, axis=0)
```

1.4 Transform data for RNN

RNN requires input of three dimensions: * The first dimension is the number of observations (games). * The second dimension is the number of time steps (minutes). * The third dimension is the number of features (in this case, 1, which is the score difference).

Our input variable (X) is a sequence of 36 observations. We will use the **reshape** method to reshape the data into a 3D array. The **reshape** method will return a 3D array with the shape (861, 36, 1) for the training data and (369, 36, 1) for the test data.

```
[ ]: X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
      X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))

      print(X_train.shape)
      print(X_test.shape)
```

```
(861, 36, 1)
(369, 36, 1)
```

1.5 MODEL02 - Simple RNN with one layer

```
[ ]: n_steps = 36
      n_inputs = 1

      model = keras.models.Sequential(
          [
              keras.layers.SimpleRNN(40, input_shape=[n_steps, n_inputs]),
              keras.layers.Dense(1, activation='sigmoid')
          ]
      )
      model.summary()
```

```
/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
```

```
Model: "sequential_1"
```

| Layer (type) | Output Shape | Param # |
|------------------------|--------------|---------|
| simple_rnn (SimpleRNN) | (None, 40) | 1,680 |
| dense_2 (Dense) | (None, 1) | 41 |

```
Total params: 1,721 (6.72 KB)
```

```
Trainable params: 1,721 (6.72 KB)
```

```
Non-trainable params: 0 (0.00 B)
```


1.5.1 How to interpret the output shape.

- First layer:
 - None is a wild-card representing the number of observations - this will be the number of observations in a batch that is fed into the network.
 - There are 32 RNN units and each produces one output. Therefore, 32 is the output shape.
- In the second layer, there is one unit and it produces one output. Therefore, $1 * 1 = 1$

1.5.2 Calculate the number of paramaters for the model

```
[ ]: p = 1 # number of parameters per unit (for RNN one remember gate)
b = 1 # number of biases per unit, for RNN this is 1 (for GRU this is 3, and
↳LSTM this is 4)

#####
# For first layer...

h = 40 # number of units in the layer
i = 1 # number of inputs (if we are predicting a stock price based on previous
↳stock prices, input = 1.
# NOTE: If we are predicting a stock price based on previous stock prices and
↳the price of oil, input = 2), etc.

params = p*(h**2+i*h+b*h)
print("First layer params = ", params)

#####
# for the last layer

# This is simply 40 outputs coming from the previous later plus 1 bias term
print("Last layer params = ", 40+1)
```

First layer params = 1680

Last layer params = 41

```
[ ]: from tensorflow.keras.callbacks import EarlyStopping

# If multiclass, use "sparse_categorical_crossentropy" as the loss function
model.compile(
    loss="binary_crossentropy",
    optimizer=tf.keras.optimizers.Nadam(learning_rate=0.001),
    metrics=['accuracy']
)
```

```
[ ]: import time
start_time = time.time()
```

```

early_stop = EarlyStopping(monitor='val_loss', min_delta = 0.0, patience=5,
    ↪ verbose=1, mode='min')

history = model.fit(
    X_train,
    y_train,
    epochs=epoch_num,
    validation_data=(X_test, y_test),
    callbacks=[early_stop]
)

end_time = time.time()

```

Epoch 1/5

27/27 12s 400ms/step -

accuracy: 0.2857 - loss: 0.9109 - val_accuracy: 0.8157 - val_loss: 0.5214

Epoch 2/5

27/27 14s 516ms/step -

accuracy: 0.7808 - loss: 0.5321 - val_accuracy: 0.8103 - val_loss: 0.4215

Epoch 3/5

27/27 16s 596ms/step -

accuracy: 0.8074 - loss: 0.4473 - val_accuracy: 0.8184 - val_loss: 0.4052

Epoch 4/5

27/27 15s 557ms/step -

accuracy: 0.8231 - loss: 0.4224 - val_accuracy: 0.8049 - val_loss: 0.3975

Epoch 5/5

27/27 16s 585ms/step -

accuracy: 0.8188 - loss: 0.4115 - val_accuracy: 0.8076 - val_loss: 0.3946

```

[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))
fig.suptitle('RNN Model Evaluation')


ax1.set_title('loss')
ax1.set_ylim(0, 1)
ax1.plot(history.history['loss'])
ax1.plot(history.history['val_loss'])
ax1.legend(['train', 'test'], loc='upper left')

ax2.set_title('accuracy')
ax2.set_ylim(0, 1)
ax2.plot(history.history['accuracy'])
ax2.plot(history.history['val_accuracy'])
ax2.legend(['train', 'test'], loc='upper left')

plt.show()

```

RNN Model Evaluation



```
[ ]: from sklearn.metrics import classification_report

y_pred = model.predict(X_test) # these predictions are probabilities (0-1)
    ↪ taken from the sigmoid function
y_pred = (y_pred > 0.5) # let's turn these into predictions (0 or 1)

print(classification_report(y_test, y_pred, digits=4))
```

```
12/12          1s 68ms/step
      precision    recall  f1-score   support

         0       0.8158      0.7425      0.7774        167
         1       0.8018      0.8614      0.8305        202

   accuracy                   0.8076        369
  macro avg       0.8088      0.8020      0.8040        369
 weighted avg       0.8082      0.8076      0.8065        369
```

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↪ f1_score

df_result = pd.DataFrame(
    {
        'model': ['RNN'],
        'time': [end_time - start_time],
        'accuracy': [accuracy_score(y_test, y_pred)],
```

```

        'precision': [precision_score(y_test, y_pred)],
        'recall': [recall_score(y_test, y_pred)],
        'f1': [f1_score(y_test, y_pred)]
    }
)

df_results = pd.concat([df_results, df_result], ignore_index=True, axis=0)

```

1.6 MODEL03 - Deep RNN

Be careful: when stacking RNN layers, you have to set “return_sequences” to True. This enables the layer to send a “sequence” of values to the next layer – just like how it uses a sequence of values for training.

Since the last layer is DENSE, it can’t take sequence data. Therefore, you CANNOT return sequences from the previous layer. So, remove return_sequences from previous layer.

```

[ ]: n_steps = 36
     n_inputs = 1

model = keras.models.Sequential(
    [
        keras.layers.SimpleRNN(32, return_sequences=True, input_shape=[n_steps,
↪n_inputs]),
        keras.layers.SimpleRNN(16, return_sequences=True),
        keras.layers.SimpleRNN(32, return_sequences=False),
        keras.layers.Dense(1, activation='sigmoid')
    ]
)

```

```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

```

```

[ ]: model.compile(
      loss="binary_crossentropy",
      optimizer=keras.optimizers.Nadam(learning_rate=0.001),
      metrics=['accuracy']
    )

model.summary()

```

Model: "sequential_2"

Layer (type)

Output Shape

Param #

| | | |
|--------------------------|----------------|-------|
| simple_rnn_1 (SimpleRNN) | (None, 36, 32) | 1,088 |
| simple_rnn_2 (SimpleRNN) | (None, 36, 16) | 784 |
| simple_rnn_3 (SimpleRNN) | (None, 32) | 1,568 |
| dense_3 (Dense) | (None, 1) | 33 |

Total params: 3,473 (13.57 KB)

Trainable params: 3,473 (13.57 KB)

Non-trainable params: 0 (0.00 B)

1.6.1 How to interpret the output shape?

- First layer:
 - None is a wild-card representing the number of observations - this will be the number of observations in a batch that is fed into the network.
 - There are 36 features (which are the ‘time steps’ of values in the sequence), and there are 32 RNN units.
 - * (batch size, time steps, features) = (None, 36, 32)
- etc. (see previous models where I elaborate on the output shape - writing this is getting repetitive)

1.6.2 Calculate the number of parameters for the model

```
[ ]: p = 1 # number of parameters per unit (for RNN one remember gate)
b = 1 # number of biases per unit, for RNN this is 1 (for GRU this is 3, and
↳LSTM this is 4)

#####
# For first layer...

h = 32 # number of units in the layer
i = 1 # number of inputs (if we are predicting a stock price based on previous
↳stock prices, input = 1.
# NOTE: If we are predicting a stock price based on previous stock prices and
↳the price of oil, input = 2), etc.

params = p*(h**2+i*h+b*h)
print("First layer params = ", params)
```

```
#####
# For second layer...

h = 16 # number of units in the layer
i = 32 # number of inputs (if we are predicting a stock price based on
↳previous stock prices, input = 1.
# NOTE: If we are predicting a stock price based on previous stock prices and
↳the price of oil, input = 2), etc.

params = p*(h**2+i*h+b*h)
print("Second layer params = ", params)

#####
# For third layer...

h = 32 # number of units in the layer
i = 16 # number of inputs (previous layer has 16 inutes, so we have 16 inputs)

params = p*(h**2+i*h+b*h)
print("Third layer params = ", params)

#####
# for the last layer

# This is simply 32 outputs coming from the previous later plus 1 bias term
print("Last layer params = ", 32+1)
```

```
First layer params = 1088
Second layer params = 784
Third layer params = 1568
Last layer params = 33
```

```
[ ]: from tensorflow.keras.callbacks import EarlyStopping

# If multiclass, use "sparse_categorical_crossentropy" as the loss function
model.compile(
    loss="binary_crossentropy",
    optimizer=tf.keras.optimizers.Nadam(learning_rate=0.001),
    metrics=['accuracy']
)
```

```
[ ]: import time
start_time = time.time()

early_stop = EarlyStopping(monitor='val_loss', min_delta=0.0, patience=5,
↳verbose=1, mode='auto')
```

```

history = model.fit(
    X_train,
    y_train,
    epochs=epoch_num,
    validation_data = (X_test, y_test),
    callbacks=[early_stop]
)

end_time = time.time()

```

```

Epoch 1/5
27/27          76s 3s/step -
accuracy: 0.7556 - loss: 0.5134 - val_accuracy: 0.8211 - val_loss: 0.4012
Epoch 2/5
27/27          69s 3s/step -
accuracy: 0.8024 - loss: 0.4187 - val_accuracy: 0.8211 - val_loss: 0.3957
Epoch 3/5
27/27          74s 3s/step -
accuracy: 0.8055 - loss: 0.4055 - val_accuracy: 0.8184 - val_loss: 0.3937
Epoch 4/5
27/27          74s 3s/step -
accuracy: 0.8116 - loss: 0.3973 - val_accuracy: 0.8157 - val_loss: 0.3933
Epoch 5/5
27/27          73s 3s/step -
accuracy: 0.8119 - loss: 0.3909 - val_accuracy: 0.8157 - val_loss: 0.3930

```

```

[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))
fig.suptitle('Deep RNN Model Evaluation')


ax1.set_title('loss')
ax1.set_ylim(0, 1)
ax1.plot(history.history['loss'])
ax1.plot(history.history['val_loss'])
ax1.legend(['train', 'test'], loc='upper left')

ax2.set_title('accuracy')
ax2.set_ylim(0, 1)
ax2.plot(history.history['accuracy'])
ax2.plot(history.history['val_accuracy'])
ax2.legend(['train', 'test'], loc='upper left')

plt.show()

```

Deep RNN Model Evaluation



```
[ ]: from sklearn.metrics import classification_report

y_pred = model.predict(X_test) # these predictions are probabilities (0-1)
    ↪ taken from the sigmoid function
y_pred = (y_pred > 0.5) # let's turn these into predictions (0 or 1)

print(classification_report(y_test, y_pred, digits=4))
```

```
12/12          4s 329ms/step
      precision    recall  f1-score   support

         0       0.8113       0.7725       0.7914         167
         1       0.8190       0.8515       0.8350         202

 accuracy                   0.8157         369
 macro avg       0.8152       0.8120       0.8132         369
 weighted avg    0.8156       0.8157       0.8152         369
```

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↪ f1_score

df_result = pd.DataFrame(
    {
        'model': ['Deep RNN'],
        'time': [end_time - start_time],
        'accuracy': [accuracy_score(y_test, y_pred)],
```



```

        'precision': [precision_score(y_test, y_pred)],
        'recall': [recall_score(y_test, y_pred)],
        'f1': [f1_score(y_test, y_pred)]
    }
)

df_results = pd.concat([df_results, df_result], ignore_index=True, axis=0)

```

1.7 MODEL04 - LSTM with one layer

```

[ ]: n_steps = 36
     n_inputs = 1

     model = keras.models.Sequential(
         [
             keras.layers.LSTM(32, input_shape=[n_steps, n_inputs]),
             keras.layers.Dense(1, activation='sigmoid')
         ]
     )

```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

super().__init__(**kwargs)

```

```

[ ]: model.compile(
     loss="binary_crossentropy",
     optimizer=keras.optimizers.Nadam(learning_rate=0.001),
     metrics=['accuracy']
)
model.summary()

```

Model: "sequential_3"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| lstm (LSTM) | (None, 32) | 4,352 |
| dense_4 (Dense) | (None, 1) | 33 |

Total params: 4,385 (17.13 KB)

Trainable params: 4,385 (17.13 KB)

Non-trainable params: 0 (0.00 B)

1.7.1 Calculate the number of paramaters for the model

```
[ ]: p = 4 # number of parameters per unit (for LSTM: 1 forget gate, 2 input gates,
      ↪and 1 output gate)
      b = 1 # number of biases per unit, for LSTM this is 1

      #####
      # For first layer....

      h = 32 # number of units in the layer
      i = 1 # number of inputs (if we are predicting a stock price based on previous
      ↪stock prices, input = 1.
      # NOTE: If we are predicting a stock price based on previous stock prices and
      ↪the price of oil, input = 2), etc.

      params = p*(h**2+i*h+b*h)
      print("First layer params = ", params)

      #####
      # for the last layer

      # This is simply 32 outputs coming from the previous later plus 1 bias term
      print("Last layer params = ", 32+1)
```

First layer params = 4352

Last layer params = 33

```
[ ]: import time
      start_time = time.time()

      early_stop = EarlyStopping(monitor='val_loss', min_delta = 0.0, patience=5,
      ↪verbose=1, mode='min')

      history = model.fit(
          X_train,
          y_train,
          epochs=epoch_num,
          validation_data=(X_test, y_test),
          callbacks=[early_stop]
      )

      end_time = time.time()
```

Epoch 1/5

27/27

3s 48ms/step -

```

accuracy: 0.4131 - loss: 0.7229 - val_accuracy: 0.8103 - val_loss: 0.5309
Epoch 2/5
27/27          1s 24ms/step -
accuracy: 0.7732 - loss: 0.5184 - val_accuracy: 0.8320 - val_loss: 0.4410
Epoch 3/5
27/27          1s 24ms/step -
accuracy: 0.7906 - loss: 0.4602 - val_accuracy: 0.8293 - val_loss: 0.4148
Epoch 4/5
27/27          1s 23ms/step -
accuracy: 0.8086 - loss: 0.4378 - val_accuracy: 0.8320 - val_loss: 0.4026
Epoch 5/5
27/27          1s 21ms/step -
accuracy: 0.8094 - loss: 0.4276 - val_accuracy: 0.8293 - val_loss: 0.3988

```

```

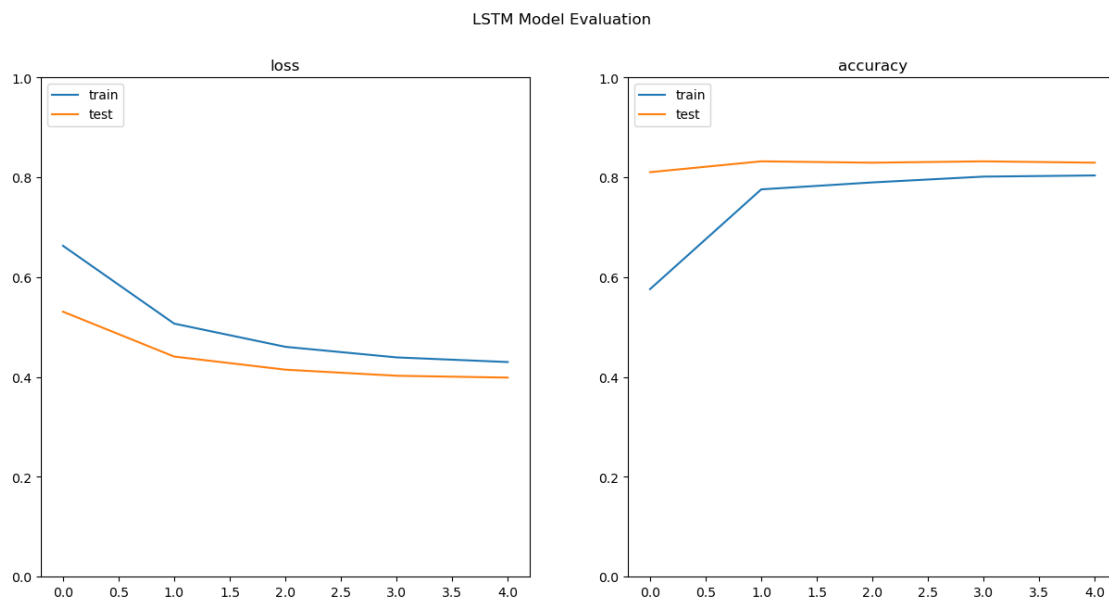
[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))
fig.suptitle('LSTM Model Evaluation')

ax1.set_title('loss')
ax1.set_ylim(0, 1)
ax1.plot(history.history['loss'])
ax1.plot(history.history['val_loss'])
ax1.legend(['train', 'test'], loc='upper left')

ax2.set_title('accuracy')
ax2.set_ylim(0, 1)
ax2.plot(history.history['accuracy'])
ax2.plot(history.history['val_accuracy'])
ax2.legend(['train', 'test'], loc='upper left')

plt.show()

```



```
[ ]: from sklearn.metrics import classification_report

y_pred = model.predict(X_test) # these predictions are probabilities (0-1)
    ↳ taken from the sigmoid function
y_pred = (y_pred > 0.5) # let's turn these into predictions (0 or 1)

print(classification_report(y_test, y_pred, digits=4))
```

```
12/12          0s 15ms/step
      precision    recall  f1-score   support

     0       0.8210      0.7964      0.8085        167
     1       0.8357      0.8564      0.8460        202

 accuracy                   0.8293        369
 macro avg       0.8284      0.8264      0.8272        369
weighted avg       0.8291      0.8293      0.8290        369
```

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳ f1_score

df_result = pd.DataFrame(
    {
        'model': ['LSTM'],
        'time': [end_time - start_time],
        'accuracy': [accuracy_score(y_test, y_pred)],
        'precision': [precision_score(y_test, y_pred)],
        'recall': [recall_score(y_test, y_pred)],
        'f1': [f1_score(y_test, y_pred)]
    }
)

df_results = pd.concat([df_results, df_result], ignore_index=True, axis=0)
```

1.8 MODEL05 - LSTM with more layers

```
[ ]: n_steps = 36
     n_inputs = 1

model = keras.models.Sequential([
    keras.layers.LSTM(32, return_sequences=True, input_shape=[n_steps,
    ↳ n_inputs]),
    keras.layers.LSTM(32, return_sequences=True),
    keras.layers.LSTM(32),
```

```
keras.layers.Dense(1, activation='sigmoid')
])
```

```
/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
```

```
[ ]: np.random.seed(42)
      tf.random.set_seed(42)

      model.compile(
          loss="binary_crossentropy",
          optimizer=keras.optimizers.Nadam(learning_rate=0.001),
          metrics=['accuracy']
      )

      model.summary()
```

Model: "sequential_4"

| Layer (type) | Output Shape | Param # |
|-----------------|----------------|---------|
| lstm_1 (LSTM) | (None, 36, 32) | 4,352 |
| lstm_2 (LSTM) | (None, 36, 32) | 8,320 |
| lstm_3 (LSTM) | (None, 32) | 8,320 |
| dense_5 (Dense) | (None, 1) | 33 |

Total params: 21,025 (82.13 KB)

Trainable params: 21,025 (82.13 KB)

Non-trainable params: 0 (0.00 B)

1.8.1 Calculate the number of paramaters for the model

```
[ ]: # Everything is the same as calculating the RNN, except for p
p = 4 # number of parameters per unit (for LSTM - 1 forget gate, 2 input gates,
    ↪ and 1 output gate)
b = 1 # number of biases per unit, for LSTM this is 1

#####
# For first layer....

h = 32 # number of units in the layer
i = 1 # number of inputs (if we are predicting a stock price based on previous
    ↪ stock prices, input = 1.
# NOTE: If we are predicting a stock price based on previous stock prices and
    ↪ the price of oil, input = 2), etc.

params = p*(h**2+i*h+b*h)
print("First layer params = ", params)

#####
# for the second layer

h = 32 # number of units in the layer
i = 32 # number of inputs (previous layer has 32 units, therefore 32 incoming
    ↪ inputs)

params = p*(h**2+i*h+b*h)
print("Second layer params = ", params)

#####
# for the third layer

h = 32 # number of units in the layer
i = 32 # number of inputs (previous layer has 32 units, therefore 32 incoming
    ↪ inputs)

params = p*(h**2+i*h+b*h)
print("Third layer params = ", params)

#####
# for the last layer

# This is simply 32 outputs coming from the previous later plus 1 bias term
print("Last layer params = ", 32+1)
```

First layer params = 4352

Second layer params = 8320

Third layer params = 8320
Last layer params = 33

```
[ ]: import time
start_time = time.time()

early_stop = EarlyStopping(monitor='val_loss', min_delta = 0.0, patience=5,
    verbose=1, mode='min')

history = model.fit(
    X_train,
    y_train,
    epochs=epoch_num,
    validation_data=(X_test, y_test),
    callbacks=[early_stop] # we defined this early_stop callback function
    earlier
)

end_time = time.time()
```

Epoch 1/5

27/27 7s 84ms/step -
accuracy: 0.7008 - loss: 0.6011 - val_accuracy: 0.7724 - val_loss: 0.4792

Epoch 2/5

27/27 1s 43ms/step -
accuracy: 0.7459 - loss: 0.5042 - val_accuracy: 0.7995 - val_loss: 0.4370

Epoch 3/5

27/27 1s 41ms/step -
accuracy: 0.7561 - loss: 0.4732 - val_accuracy: 0.8211 - val_loss: 0.4217

Epoch 4/5

27/27 1s 42ms/step -
accuracy: 0.7630 - loss: 0.4568 - val_accuracy: 0.8184 - val_loss: 0.4151

Epoch 5/5

27/27 1s 41ms/step -
accuracy: 0.7813 - loss: 0.4445 - val_accuracy: 0.8184 - val_loss: 0.4112

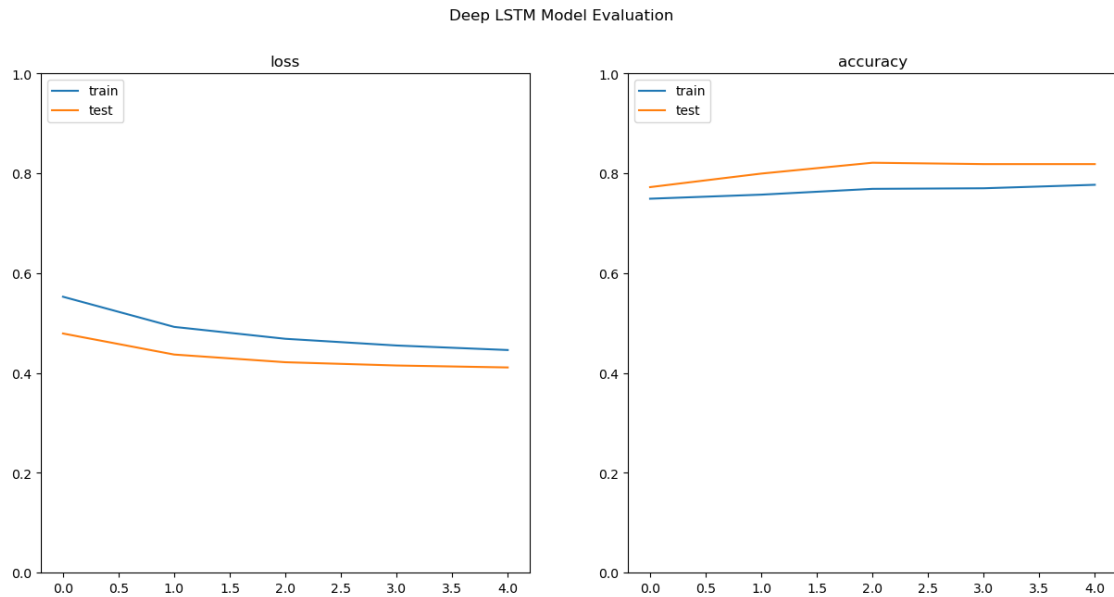
```
[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))
fig.suptitle('Deep LSTM Model Evaluation')

ax1.set_title('loss')
ax1.set_ylim(0, 1)
ax1.plot(history.history['loss'])
ax1.plot(history.history['val_loss'])
ax1.legend(['train', 'test'], loc='upper left')

ax2.set_title('accuracy')
ax2.set_ylim(0, 1)
ax2.plot(history.history['accuracy'])
```

```
ax2.plot(history.history['val_accuracy'])
ax2.legend(['train', 'test'], loc='upper left')

plt.show()
```



```
[ ]: from sklearn.metrics import classification_report

y_pred = model.predict(X_test) # these predictions are probabilities (0-1)
    ↳ taken from the sigmoid function
y_pred = (y_pred > 0.5) # let's turn these into predictions (0 or 1)

print(classification_report(y_test, y_pred, digits=4))
```

```
12/12          1s 47ms/step
      precision    recall  f1-score   support

     0       0.8425     0.7365     0.7859        167
     1       0.8027     0.8861     0.8424        202

 accuracy                   0.8184        369
 macro avg       0.8226     0.8113     0.8141        369
 weighted avg    0.8207     0.8184     0.8168        369
```

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳ f1_score
```



```

df_result = pd.DataFrame(
    {
        'model': ['Deep LSTM'],
        'time': [end_time - start_time],
        'accuracy': [accuracy_score(y_test, y_pred)],
        'precision': [precision_score(y_test, y_pred)],
        'recall': [recall_score(y_test, y_pred)],
        'f1': [f1_score(y_test, y_pred)]
    }
)

df_results = pd.concat([df_results, df_result], ignore_index=True, axis=0)

```

1.9 MODEL06 - GRU with one layer

```

[ ]: n_steps = 36
     n_inputs = 1

     model = keras.models.Sequential([
         keras.layers.GRU(32, input_shape=[n_steps, n_inputs]),
         keras.layers.Dense(1, activation='sigmoid')
     ])

```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(**kwargs)

```

[ ]: model.compile(
     loss="binary_crossentropy",
     optimizer=keras.optimizers.Nadam(learning_rate=0.001),
     metrics=['accuracy']
)
model.summary()

```

Model: "sequential_5"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| gru (GRU) | (None, 32) | 3,360 |
| dense_6 (Dense) | (None, 1) | 33 |

Total params: 3,393 (13.25 KB)

Trainable params: 3,393 (13.25 KB)

Non-trainable params: 0 (0.00 B)

1.9.1 Calculate the number of paramaters for the model

```
[ ]: # Everything is the same as calculating the RNN or LSTM parameters, except for p
p = 3 # number of parameters per unit (for GRU - 1 forget gate, 1 input gates,
    ↪and 1 output gate)
b = 2 # number of biases per unit, for GRU this is 2

#####
# For first layer....

h = 32 # number of units in the layer
i = 1 # number of inputs (if we are predicting a stock price based on previous
    ↪stock prices, input = 1.
# NOTE: If we are predicting a stock price based on previous stock prices and
    ↪the price of oil, input = 2), etc.

params = p*(h**2+i*h+b*h)
print("First layer params = ", params)

#####
# for the last layer

# This is simply 32 outputs coming from the previous later plus 1 bias term
print("Last layer params = ", 32+1)
```

First layer params = 3360

Last layer params = 33

```
[ ]: import time
start_time = time.time()

early_stop = EarlyStopping(monitor='val_loss', min_delta = 0.0, patience=5,
    ↪verbose=1, mode='min')


history = model.fit(
    X_train,
    y_train,
    epochs=epoch_num,
    validation_data=(X_test, y_test),
    callbacks=[early_stop] # we defined this early_stop callback function
    ↪earlier
```

```
)  
  
end_time = time.time()
```

```
Epoch 1/5  
27/27          3s 49ms/step -  
accuracy: 0.7795 - loss: 0.4419 - val_accuracy: 0.8266 - val_loss: 0.3876  
Epoch 2/5  
27/27          1s 23ms/step -  
accuracy: 0.7829 - loss: 0.4220 - val_accuracy: 0.8184 - val_loss: 0.3853  
Epoch 3/5  
27/27          1s 24ms/step -  
accuracy: 0.7908 - loss: 0.4133 - val_accuracy: 0.8130 - val_loss: 0.3851  
Epoch 4/5  
27/27          1s 22ms/step -  
accuracy: 0.7955 - loss: 0.4095 - val_accuracy: 0.8184 - val_loss: 0.3848  
Epoch 5/5  
27/27          1s 21ms/step -  
accuracy: 0.7946 - loss: 0.4074 - val_accuracy: 0.8211 - val_loss: 0.3845
```

```
[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))  
fig.suptitle('GRU Model Evaluation')  
  
ax1.set_title('loss')  
ax1.set_ylim(0, 1)  
ax1.plot(history.history['loss'])  
ax1.plot(history.history['val_loss'])  
ax1.legend(['train', 'test'], loc='upper left')  
  
ax2.set_title('accuracy')  
ax2.set_ylim(0, 1)  
ax2.plot(history.history['accuracy'])  
ax2.plot(history.history['val_accuracy'])  
ax2.legend(['train', 'test'], loc='upper left')  
  
plt.show()
```

GRU Model Evaluation



```
[ ]: from sklearn.metrics import classification_report

y_pred = model.predict(X_test) # these predictions are probabilities (0-1)
    ↳ taken from the sigmoid function
y_pred = (y_pred > 0.5) # let's turn these into predictions (0 or 1)

print(classification_report(y_test, y_pred, digits=4))
```

```
12/12          1s 19ms/step
      precision    recall  f1-score   support

         0       0.8176    0.7784    0.7975        167
         1       0.8238    0.8564    0.8398        202

   accuracy                   0.8211        369
  macro avg       0.8207    0.8174    0.8187        369
weighted avg       0.8210    0.8211    0.8207        369
```

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳ f1_score

df_result = pd.DataFrame(
    {
        'model': ['GRU'],
        'time': [end_time - start_time],
        'accuracy': [accuracy_score(y_test, y_pred)],
```

```

        'precision': [precision_score(y_test, y_pred)],
        'recall': [recall_score(y_test, y_pred)],
        'f1': [f1_score(y_test, y_pred)]
    }
)

df_results = pd.concat([df_results, df_result], ignore_index=True, axis=0)

```

1.10 MODEL07 - GRU with more layers

```

[ ]: n_steps = 36
     n_inputs = 1

model = keras.models.Sequential([
    keras.layers.GRU(32, return_sequences=True, input_shape=[n_steps,
↪n_inputs]),
    keras.layers.GRU(32, return_sequences=True),
    keras.layers.GRU(32),
    keras.layers.Dense(1, activation='sigmoid')
])

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

```

```

[ ]: model.compile(
      loss="binary_crossentropy",
      optimizer=keras.optimizers.Nadam(learning_rate=0.001),
      metrics=['accuracy']
    )
model.summary()

```

Model: "sequential_6"

| Layer (type) | Output Shape | Param # |
|-----------------|----------------|---------|
| gru_1 (GRU) | (None, 36, 32) | 3,360 |
| gru_2 (GRU) | (None, 36, 32) | 6,336 |
| gru_3 (GRU) | (None, 32) | 6,336 |
| dense_7 (Dense) | (None, 1) | 33 |

Total params: 16,065 (62.75 KB)

Trainable params: 16,065 (62.75 KB)

Non-trainable params: 0 (0.00 B)

1.10.1 Calculate the number of paramaters for the model

```
[ ]: # Everything is the same as calculating the RNN or LSTM parameters, except for p
p = 3 # number of parameters per unit (for GRU - 1 forget gate, 1 input gates,
      ↪and 1 output gate)
b = 2 # number of biases per unit, for GRU this is 2

#####
# For first layer....

h = 32 # number of units in the layer
i = 1  # number of inputs (if we are predicting a stock price based on previous
      ↪stock prices, input = 1.
# NOTE: If we are predicting a stock price based on previous stock prices and
      ↪the price of oil, input = 2), etc.

params = p*(h**2+i*h+b*h)
print("First layer params = ", params)

#####
# For second layer....

h = 32 # number of units in the layer
i = 32 # number of inputs (previous layer has 32 units, therefore 32 inputs)

params = p*(h**2+i*h+b*h)
print("First layer params = ", params)

#####
# For third layer....

h = 32 # number of units in the layer
i = 32 # number of inputs (previous layer has 32 units, therefore 32 inputs)

params = p*(h**2+i*h+b*h)
print("First layer params = ", params)

#####
```

```
# for the last layer

# This is simply 32 outputs coming from the previous later plus 1 bias term
print("Last layer params = ", 32+1)
```

```
First layer params = 3360
First layer params = 6336
First layer params = 6336
Last layer params = 33
```

```
[ ]: import time
start_time = time.time()

early_stop = EarlyStopping(monitor='val_loss', min_delta = 0.0, patience=5,
    ↳ verbose=1, mode='min')

history = model.fit(
    X_train,
    y_train,
    epochs=epoch_num,
    validation_data=(X_test, y_test),
    callbacks=[early_stop]
)

end_time = time.time()
```

```
Epoch 1/5
27/27          6s 82ms/step -
accuracy: 0.5762 - loss: 0.6405 - val_accuracy: 0.8320 - val_loss: 0.4160
Epoch 2/5
27/27          1s 41ms/step -
accuracy: 0.7628 - loss: 0.4734 - val_accuracy: 0.8293 - val_loss: 0.4025
Epoch 3/5
27/27          1s 40ms/step -
accuracy: 0.7727 - loss: 0.4435 - val_accuracy: 0.8347 - val_loss: 0.3909
Epoch 4/5
27/27          1s 40ms/step -
accuracy: 0.7716 - loss: 0.4266 - val_accuracy: 0.8320 - val_loss: 0.3868
Epoch 5/5
27/27          1s 40ms/step -
accuracy: 0.7776 - loss: 0.4176 - val_accuracy: 0.8238 - val_loss: 0.3840
```

```
[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))
fig.suptitle('Deep GRU Model Evaluation')

ax1.set_title('loss')
ax1.set_ylim(0, 1)
ax1.plot(history.history['loss'])
```


```

ax1.plot(history.history['val_loss'])
ax1.legend(['train', 'test'], loc='upper left')

ax2.set_title('accuracy')
ax2.set_ylim(0, 1)
ax2.plot(history.history['accuracy'])
ax2.plot(history.history['val_accuracy'])
ax2.legend(['train', 'test'], loc='upper left')

plt.show()

```



```

[ ]: from sklearn.metrics import classification_report

y_pred = model.predict(X_test) # these predictions are probabilities (0-1)
    ↪ taken from the sigmoid function
y_pred = (y_pred > 0.5) # let's turn these into predictions (0 or 1)

print(classification_report(y_test, y_pred, digits=4))

```

```

12/12          1s 34ms/step
          precision    recall  f1-score   support

         0       0.8228      0.7784      0.8000        167
         1       0.8246      0.8614      0.8426        202

 accuracy                   0.8238        369
 macro avg       0.8237      0.8199      0.8213        369
 weighted avg     0.8238      0.8238      0.8233        369

```



```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

df_result = pd.DataFrame(
    {
        'model': ['Deep GRU'],
        'time': [end_time - start_time],
        'accuracy': [accuracy_score(y_test, y_pred)],
        'precision': [precision_score(y_test, y_pred)],
        'recall': [recall_score(y_test, y_pred)],
        'f1': [f1_score(y_test, y_pred)]
    }
)

df_results = pd.concat([df_results, df_result], ignore_index=True, axis=0)
```

1.11 MODEL08 - LSTM with Conv1D

```
[ ]: n_steps = 36
n_inputs = 1

model = keras.models.Sequential([
    keras.layers.Conv1D(filters=10, kernel_size=3, strides=1, padding="valid",
        input_shape=[n_steps, n_inputs]),
    keras.layers.Conv1D(filters=20, kernel_size=3, strides=1, padding="valid"),
    keras.layers.LSTM(32, return_sequences=True),
    keras.layers.LSTM(32),
    keras.layers.Dense(1, activation='sigmoid')
])

model.summary()
```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-packages/keras/src/layers/convolutional/base_conv.py:99: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(
Model: "sequential_7"
```

| Layer (type) | Output Shape | Param # |
|-----------------|----------------|---------|
| conv1d (Conv1D) | (None, 34, 10) | 40 |

| | | |
|-------------------|----------------|-------|
| conv1d_1 (Conv1D) | (None, 32, 20) | 620 |
| lstm_4 (LSTM) | (None, 32, 32) | 6,784 |
| lstm_5 (LSTM) | (None, 32) | 8,320 |
| dense_8 (Dense) | (None, 1) | 33 |

Total params: 15,797 (61.71 KB)

Trainable params: 15,797 (61.71 KB)

Non-trainable params: 0 (0.00 B)

```
[ ]: model.compile(
    loss="binary_crossentropy",
    optimizer=keras.optimizers.Nadam(learning_rate=0.010),
    metrics=['accuracy']
)
```

```
[ ]: import time
start_time = time.time()

early_stop = EarlyStopping(monitor='val_loss', min_delta = 0.0, patience=5,
    verbose=1, mode='min')

history = model.fit(
    X_train,
    y_train,
    epochs=epoch_num,
    validation_data=(X_test, y_test),
    callbacks=[early_stop]
)

end_time = time.time()
```

Epoch 1/5

27/27 8s 91ms/step -

accuracy: 0.7545 - loss: 0.5260 - val_accuracy: 0.8157 - val_loss: 0.3957

Epoch 2/5

27/27 1s 44ms/step -

accuracy: 0.7831 - loss: 0.4275 - val_accuracy: 0.8103 - val_loss: 0.3928

Epoch 3/5

27/27 1s 40ms/step -

accuracy: 0.7988 - loss: 0.4136 - val_accuracy: 0.8184 - val_loss: 0.3947

Epoch 4/5

27/27 1s 40ms/step -

accuracy: 0.8078 - loss: 0.4101 - val_accuracy: 0.8238 - val_loss: 0.4010

Epoch 5/5

27/27 1s 41ms/step -

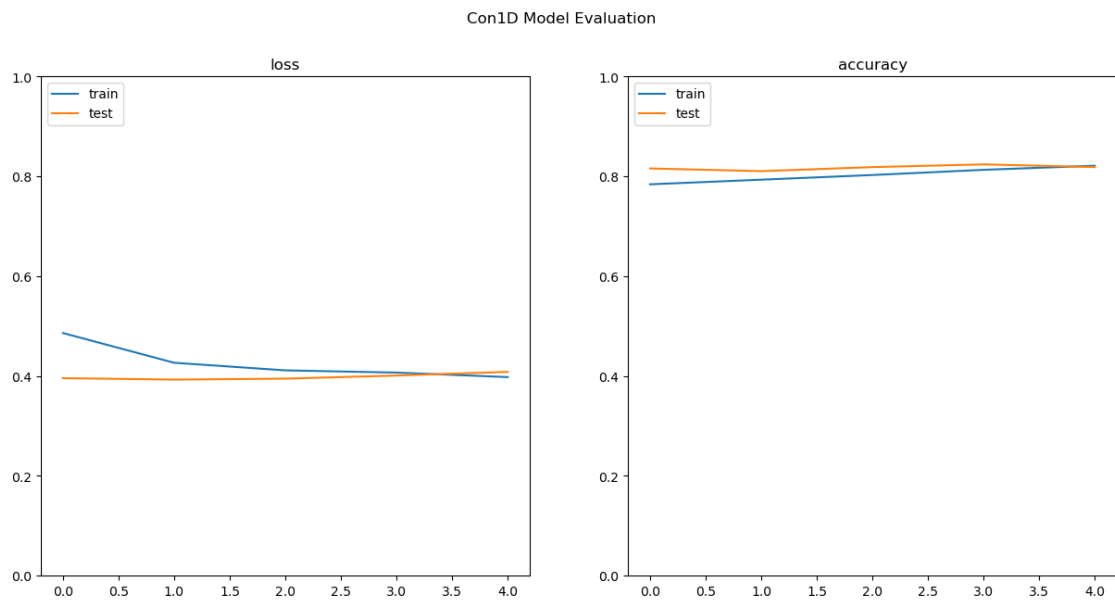
accuracy: 0.8257 - loss: 0.4012 - val_accuracy: 0.8184 - val_loss: 0.4084

```
[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 7))
fig.suptitle('Con1D Model Evaluation')

ax1.set_title('loss')
ax1.set_ylim(0, 1)
ax1.plot(history.history['loss'])
ax1.plot(history.history['val_loss'])
ax1.legend(['train', 'test'], loc='upper left')

ax2.set_title('accuracy')
ax2.set_ylim(0, 1)
ax2.plot(history.history['accuracy'])
ax2.plot(history.history['val_accuracy'])
ax2.legend(['train', 'test'], loc='upper left')

plt.show()
```



```
[ ]: from sklearn.metrics import classification_report
```

```

y_pred = model.predict(X_test) # these predictions are probabilities (0-1)
    ↳ taken from the sigmoid function
y_pred = (y_pred > 0.5) # let's turn these into predictions (0 or 1)

print(classification_report(y_test, y_pred, digits=4))

```

```

12/12          1s 42ms/step
      precision    recall  f1-score   support

     0       0.8205       0.7665       0.7926         167
     1       0.8169       0.8614       0.8386         202

 accuracy                   0.8184         369
 macro avg       0.8187       0.8139       0.8156         369
 weighted avg    0.8185       0.8184       0.8177         369

```

```

[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳ f1_score

df_result = pd.DataFrame(
    {
        'model': ['Conv1D Deep LSTM'],
        'time': [end_time - start_time],
        'accuracy': [accuracy_score(y_test, y_pred)],
        'precision': [precision_score(y_test, y_pred)],
        'recall': [recall_score(y_test, y_pred)],
        'f1': [f1_score(y_test, y_pred)]
    }
)

df_results = pd.concat([df_results, df_result], ignore_index=True, axis=0)

```

1.12 Summary of results

```

[ ]: df_results.sort_values(by=['accuracy', 'precision', 'recall', 'f1'],
    ↳ ascending=False)

```

```

[ ]:
      model      time  accuracy  precision  recall      f1
3      LSTM   5.880496   0.829268   0.835749   0.856436   0.845966
6    Deep GRU  10.679676   0.823848   0.824645   0.861386   0.842615
5      GRU    5.929280   0.821138   0.823810   0.856436   0.839806
7 Conv1D Deep LSTM  12.128046   0.818428   0.816901   0.861386   0.838554
4    Deep LSTM  11.437630   0.818428   0.802691   0.886139   0.842353
2    Deep RNN  367.517398   0.815718   0.819048   0.851485   0.834951
1      RNN    72.692728   0.807588   0.801843   0.861386   0.830549
0      NN     3.624414   0.794038   0.850000   0.757426   0.801047

```