

# Tutorial2-RNN Time series-ALL DATA Standardization

April 9, 2024

## 1 Tutorial 2 - RNN Time Series

In this notebook, we will predict the weather temperature.

```
[ ]: import tensorflow as tf
      from tensorflow import keras
      from sklearn.metrics import mean_squared_error

      epoch_num = 5 # number of epochs to use for training our models.

      # Common imports
      import numpy as np
      import os

      # to make this notebook's output stable across runs
      np.random.seed(42)

      # To plot pretty figures
      %matplotlib inline
      import matplotlib as mpl
      import matplotlib.pyplot as plt
      mpl.rc('axes', labelsizes=14)
      mpl.rc('xtick', labelsizes=12)
      mpl.rc('ytick', labelsizes=12)
```

## 2 Read the Dataset

```
[ ]: import pandas as pd

      weather = pd.read_csv('https://raw.githubusercontent.com/prof-tcsmith/data/
      ↪master/weather.csv')

      weather.head()
```

```
[ ]:      date  hour    NO2      CO      O3      NO  PM2.5    PM10  \
0  4/25/2021     0  0.039817  0.080700 -0.000867  0.009800    0.2  5.040000
1  4/25/2021     1  0.035900  0.092217 -0.000267  0.009833    0.2  6.293333
```

```

2  4/25/2021      2  0.028083  0.062750  0.002517  0.012883  0.2  5.501667
3  4/25/2021      3  0.025633  0.042300  0.004550  0.014233  0.2  4.201667
4  4/25/2021      4  0.023717  0.036883  0.006267  0.015417  0.2  5.365000

```

```

      Air Temp.  Air Hum.    Air Pres.
0  25.795000    99.9  1011.980000
1  25.445000    99.9  1012.131667
2  25.223333    99.9  1012.365000
3  25.075000    99.9  1012.276667
4  24.928333    99.9  1012.030000

```

```
[ ]: # Convert the temp to Fahrenheit:
```

```
weather['Air Temp F'] = weather['Air Temp.']*1.8 + 32
```

```
[ ]: weather
```

```

[ ]:
      date  hour      NO2      CO      O3      NO  PM2.5  \
0    4/25/2021    0  0.039817  0.080700 -0.000867  0.009800  0.200
1    4/25/2021    1  0.035900  0.092217 -0.000267  0.009833  0.200
2    4/25/2021    2  0.028083  0.062750  0.002517  0.012883  0.200
3    4/25/2021    3  0.025633  0.042300  0.004550  0.014233  0.200
4    4/25/2021    4  0.023717  0.036883  0.006267  0.015417  0.200
...
7952  4/19/2022   19  0.007500  0.135750  0.035250  0.044000  2.150
7953  4/19/2022   20  0.024000  0.145750  0.025250  0.035000  2.525
7954  4/19/2022   21  0.013400  0.147200  0.038800  0.022200  2.260
7955  4/19/2022   22  0.023000  0.126000  0.039000  0.014000  2.350
7956  4/19/2022   23  0.032500  0.162250  0.041750  0.030500  1.975

```

```

      PM10  Air Temp.  Air Hum.    Air Pres.  Air Temp F
0    5.040000  25.795000    99.900  1011.980000    78.431
1    6.293333  25.445000    99.900  1012.131667    77.801
2    5.501667  25.223333    99.900  1012.365000    77.402
3    4.201667  25.075000    99.900  1012.276667    77.135
4    5.365000  24.928333    99.900  1012.030000    76.871
...
7952   9.275000  28.375000    39.175  1016.000000    83.075
7953  16.100000  28.575000    37.275  1015.550000    83.435
7954  12.980000  28.440000    32.740  1015.340000    83.192
7955  11.375000  28.150000    33.250  1015.600000    82.670
7956  10.675000  25.700000    41.725  1015.725000    78.260

```

```
[7957 rows x 12 columns]
```

```
[ ]: #Drop the columns we don't need
```

```
weather = weather.drop(['NO2', 'CO', 'O3', 'NO', 'PM2.5', 'PM10', 'Air Temp.',
                        'Air Hum.', 'Air Pres.'], axis=1)
```


```
[ ]: weather
```

```
[ ]:
      date  hour  Air Temp F
0   4/25/2021    0    78.431
1   4/25/2021    1    77.801
2   4/25/2021    2    77.402
3   4/25/2021    3    77.135
4   4/25/2021    4    76.871
...
7952 4/19/2022   19    83.075
7953 4/19/2022   20    83.435
7954 4/19/2022   21    83.192
7955 4/19/2022   22    82.670
7956 4/19/2022   23    78.260
```

```
[7957 rows x 3 columns]
```

```
[ ]: #Plot temp

plt.plot(weather['Air Temp F'])
plt.show()
```



### 3 Data Cleanup

```
[ ]: # Values higher than 100 degrees are probably incorrect readings
```

```
weather[weather['Air Temp F']>100]
```

```
[ ]:
      date  hour  Air Temp F
1578 6/29/2021   21    124.865
1582 6/30/2021    1    108.674
1595 6/30/2021   14    112.001
1596 6/30/2021   15    131.294
1597 6/30/2021   16    128.849
...
1742 7/6/2021   17    256.820
1743 7/6/2021   18    256.820
1744 7/6/2021   19    251.894
1745 7/6/2021   20    256.820
1746 7/6/2021   21    246.926
```

```
[114 rows x 3 columns]
```


```
[ ]: # Convert all values higher than 100 degrees to null values
```

```
weather['Air Temp F'] = np.where(weather['Air Temp F']>100, np.nan,
↪weather['Air Temp F'])
```

```
[ ]: # Values lower than 30 degrees are probably incorrect readings. Convert them to
↪null
```


```
weather['Air Temp F'] = np.where(weather['Air Temp F']<30, np.nan, weather['Air
↪Temp F'])
```

```
[ ]: plt.plot(weather['Air Temp F'])
plt.show()
```



```
[ ]: # Remove all null values  
  
weather = weather.dropna().reset_index(drop=True)
```

```
[ ]: plt.plot(weather['Air Temp F'])  
plt.show()
```



### 3.1 RESHAPE the data set!

```
[ ]: weather.shape
```

```
[ ]: (7017, 3)
```

```
[ ]: # Note that not all days have 24 readings. Some are missing.
```

```
weather.shape[0]/24
```

```
[ ]: 292.375
```

```
[ ]: weather.groupby(['date']).count()
```

```
[ ]:
      hour  Air Temp F
date
1/1/2022   24         24
1/10/2022  24         24
1/11/2022  24         24
1/12/2022  24         24
1/13/2022  24         24
...      ...      ...
```

9/5/2021	24	24
9/6/2021	24	24
9/7/2021	24	24
9/8/2021	24	24
9/9/2021	24	24

[299 rows x 2 columns]

```
[ ]: # Find the reading count for each day

hour_count = pd.DataFrame(weather.groupby(['date']).count()['hour'])

hour_count
```

```
[ ]:      hour
date
1/1/2022    24
1/10/2022   24
1/11/2022   24
1/12/2022   24
1/13/2022   24
...
9/5/2021    24
9/6/2021    24
9/7/2021    24
9/8/2021    24
9/9/2021    24
```

[299 rows x 1 columns]

```
[ ]: # Find the reading counts that are less than 24

hour_count[hour_count['hour']<24]
```

```
[ ]:      hour
date
3/18/2022    7
4/11/2022    9
4/14/2022   10
6/21/2021   21
6/29/2021   23
6/30/2021   16
7/1/2021     1
7/2/2021     3
7/3/2021     3
7/4/2021     3
8/10/2021    9
```

```
[ ]: # Identify the dates of these records

hour_count[hour_count['hour']<24].index.values
```

```
[ ]: array(['3/18/2022', '4/11/2022', '4/14/2022', '6/21/2021', '6/29/2021',
        '6/30/2021', '7/1/2021', '7/2/2021', '7/3/2021', '7/4/2021',
        '8/10/2021'], dtype=object)
```

```
[ ]: # Find the corresponding index values in the original data set

indexes = weather[weather['date'].isin(hour_count[hour_count['hour']<24].index.
↪values)]

indexes
```

```
[ ]:
      date  hour  Air Temp F
1368 6/21/2021    0    84.386
1369 6/21/2021    1    84.137
1370 6/21/2021    2    83.993
1371 6/21/2021    3    83.549
1372 6/21/2021    4    83.210
...     ...    ...
6892 4/14/2022   19    83.732
6893 4/14/2022   20    82.715
6894 4/14/2022   21    81.590
6895 4/14/2022   22    82.085
6896 4/14/2022   23    78.620

[105 rows x 3 columns]
```

```
[ ]: # Remove these rows from the data set.

weather = weather.drop(indexes.index, axis=0).reset_index(drop=True)

weather.shape
```

```
[ ]: (6912, 3)
```

```
[ ]: weather.head()
```

```
[ ]:
      date  hour  Air Temp F
0  4/25/2021    0    78.431
1  4/25/2021    1    77.801
2  4/25/2021    2    77.402
3  4/25/2021    3    77.135
4  4/25/2021    4    76.871
```




```
[ ]: # All remaining days have 24 readings (for 24 hours)
# There are a total of 288 days
```

```
weather.shape[0]/24
```

```
[ ]: 288.0
```

```
[ ]: plt.plot(weather['Air Temp F'])
plt.show()
```



```
[ ]: # Re-organize the data set by day and hours
```

```
temp = np.array(weather['Air Temp F']).reshape(288,24)
```

```
temp
```

```
[ ]: array([[78.431      , 77.801      , 77.40199999, ..., 82.54700001,
          81.716      , 79.196      ],
          [76.02200001, 73.121      , 71.68699999, ..., 84.82699999,
          84.57499999, 82.52900001],
          [80.843      , 78.87500001, 77.05099999, ..., 84.52700001,
          83.99899999, 82.44199999],
```

```
...,
[78.26      , 77.54      , 77.495      , ..., 83.98999999,
 83.3      , 79.736     ],
[78.305     , 77.63      , 77.432      , ..., 84.245      ,
 83.084     , 79.376     ],
[77.27      , 76.136     , 75.29      , ..., 83.192      ,
 82.67      , 78.26      ]])
```

```
[ ]: # Convert to dataframe
```

```
temp_df = pd.DataFrame(temp, columns=np.arange(0,24,1))
```

```
temp_df
```

```
[ ]:
```

	0	1	2	3	4	5	6	7	8	\
0	78.431	77.801	77.402	77.135	76.871	76.814	76.892	76.925	76.580	
1	76.022	73.121	71.687	70.664	69.560	68.864	68.603	68.360	68.360	
2	80.843	78.875	77.051	74.675	73.499	72.950	72.221	71.330	71.048	
3	80.576	78.731	76.739	74.819	73.829	73.052	72.575	71.876	71.306	
4	78.632	77.618	76.040	75.278	74.918	74.561	73.859	73.127	72.785	
..	...	...	...	...	...	...	...	...	...	
283	76.910	76.145	75.590	75.380	75.245	74.525	74.750	74.300	74.030	
284	75.200	77.585	77.060	76.055	74.948	74.030	73.400	72.500	71.915	
285	78.260	77.540	77.495	77.540	77.540	76.730	75.920	75.245	74.525	
286	78.305	77.630	77.432	77.135	77.360	76.640	76.505	75.980	75.065	
287	77.270	76.136	75.290	74.948	74.570	74.210	74.030	74.300	74.060	

	9	...	14	15	16	17	18	19	20	\
0	76.343	...	81.584	81.575	82.445	84.731	84.272	83.252	83.447	
1	68.267	...	75.500	77.594	79.691	81.458	83.012	84.080	84.323	
2	70.766	...	77.222	79.241	80.933	81.602	82.478	83.795	84.146	
3	70.793	...	77.657	79.562	81.014	81.848	82.955	83.813	83.996	
4	72.575	...	79.859	82.376	85.808	86.831	86.108	86.459	86.003	
..	...	...	...	...	...	...	...	...	...	
283	73.040	...	83.300	84.920	81.428	83.210	83.750	85.400	86.225	
284	71.600	...	77.108	79.916	83.012	83.930	83.948	86.000	87.215	
285	73.400	...	80.330	82.625	84.695	85.640	85.460	84.992	84.470	
286	74.120	...	80.810	83.444	84.155	83.435	80.510	86.585	85.235	
287	72.230	...	73.976	74.930	78.485	79.916	81.635	83.075	83.435	

	21	22	23
0	82.547	81.716	79.196
1	84.827	84.575	82.529
2	84.527	83.999	82.442
3	84.437	82.913	80.210
4	84.902	84.413	81.434
..	...	...	...

```
283  87.980  86.540  76.784
284  87.800  86.540  80.168
285  83.990  83.300  79.736
286  84.245  83.084  79.376
287  83.192  82.670  78.260
```

```
[288 rows x 24 columns]
```

## 4 Reshape for Standardizing Data

```
[ ]: # Let's create a single sequence (i.e., feature) for standardization
```

```
temp_1feature = np.array(temp_df).ravel().reshape(-1,1)

temp_1feature.shape
```

```
[ ]: (6912, 1)
```

```
[ ]: temp_1feature
```

```
[ ]: array([[78.431    ],
           [77.801    ],
           [77.40199999],
           ...,
           [83.192    ],
           [82.67     ],
           [78.26     ]])
```

### 4.1 Standardize the values

```
[ ]: # Next, standardize

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

temp_std = scaler.fit_transform(temp_1feature)
```

### 4.2 Reshape the data back to 24-hour format

```
[ ]: temp_resaped = temp_std.reshape(288,24)

temp_resaped.shape
```

```
[ ]: (288, 24)
```

```
[ ]: #Pandas version of the reshaped data
```

```
pd.DataFrame(temp_reshaped, columns=np.arange(0,24,1))
```

```
[ ]:
```

	0	1	2	3	4	5	6	\
0	0.391232	0.325683	0.284169	0.256388	0.228920	0.222990	0.231105	
1	0.140586	-0.161251	-0.310453	-0.416891	-0.531758	-0.604174	-0.631330	
2	0.642190	0.437428	0.247649	0.000436	-0.121922	-0.179043	-0.254892	
3	0.614410	0.422445	0.215186	0.015419	-0.087587	-0.168430	-0.218060	
4	0.412145	0.306642	0.142458	0.063175	0.025719	-0.011425	-0.084465	
..	...	...	...	...	...	...	...	
283	0.232978	0.153383	0.095638	0.073788	0.059742	-0.015171	0.008239	
284	0.055060	0.303209	0.248585	0.144019	0.028840	-0.066674	-0.132222	
285	0.373440	0.298527	0.293845	0.298527	0.298527	0.214250	0.129973	
286	0.378122	0.307891	0.287290	0.256388	0.279799	0.204886	0.190840	
287	0.270435	0.152447	0.064424	0.028840	-0.010489	-0.047945	-0.066674	
	7	8	9	...	14	15	16	\
0	0.234539	0.198643	0.173984	...	0.719288	0.718351	0.808871	
1	-0.656613	-0.656613	-0.666289	...	0.086274	0.304145	0.522329	
2	-0.347597	-0.376938	-0.406279	...	0.265440	0.475509	0.651554	
3	-0.290788	-0.350094	-0.403470	...	0.310700	0.508907	0.659982	
4	-0.160627	-0.196210	-0.218060	...	0.539809	0.801692	1.158777	
..	...	...	...	...	...	...	...	
283	-0.038581	-0.066674	-0.169679	...	0.897830	1.066384	0.703057	
284	-0.225864	-0.286730	-0.319505	...	0.253579	0.545740	0.867865	
285	0.059742	-0.015171	-0.132222	...	0.588814	0.827599	1.042974	
286	0.136216	0.041014	-0.057309	...	0.638756	0.912813	0.986789	
287	-0.038581	-0.063552	-0.253956	...	-0.072292	0.026968	0.396850	
	17	18	19	20	21	22	23	
0	1.046720	0.998963	0.892836	0.913125	0.819484	0.733022	0.470827	
1	0.706178	0.867865	0.978986	1.004269	1.056708	1.030489	0.817611	
2	0.721161	0.812305	0.949333	0.985853	1.025494	0.970558	0.808559	
3	0.746756	0.861934	0.951206	0.970246	1.016130	0.857565	0.576329	
4	1.265216	1.189991	1.226511	1.179066	1.064511	1.013633	0.703681	
..	...	...	...	...	...	...	...	
283	0.888466	0.944651	1.116326	1.202164	1.384764	1.234938	0.219868	
284	0.963379	0.965252	1.178754	1.305169	1.366036	1.234938	0.571959	
285	1.141297	1.122569	1.073876	1.019564	0.969622	0.897830	0.527011	
286	0.911876	0.607543	1.239620	1.099159	0.996153	0.875356	0.489555	
287	0.545740	0.724594	0.874420	0.911876	0.886593	0.832281	0.373440	

```
[288 rows x 24 columns]
```

## 5 Split the Data

In certain cases, we cannot use a random split. For example, if we are trying to predict the stock market, we cannot use a random split. We need to use a chronological split.

BUT, keep in mind if we have something like hourly readings of daily temperature, we can use a random split on days, but the sequence of the temperature within the day is important to remain sequential.

In this case, we are using a random split because each day as an independent sample.

```
[ ]: from sklearn.model_selection import train_test_split

train, test = train_test_split(temp_resaped, test_size=0.3)

# if we neede to maintain the ordering of the data, we can use the following ↵
↵code to split the data
#split_point = int(len(temp_resaped)*0.7)
#train, test = temp_resaped[:split_point], temp_resaped[split_point:]
```

```
[ ]: train.shape
```

```
[ ]: (201, 24)
```

```
[ ]: train[:2]
```

```
[ ]: array([[ -1.98631703, -1.8364912 , -1.85053737, -1.75689623, -1.90672206,
          -1.94886057, -1.92545029, -1.85521943, -1.79435268, -1.79435268,
          -1.76157828, -1.7709424 , -1.81776297, -1.87394766, -1.84117326,
          -1.82712708, -1.80839886, -1.74285005, -1.65857302, -1.55556776,
          -1.49001896, -1.46192662, -1.45256251, -1.42447016],
          [ 0.40153216,  0.44835273,  0.48580919,  0.47644508,  0.49985536,
           0.51858359,  0.54199388,  0.52794771,  0.54199388,  0.54667593,
           0.51858359,  0.57476828,  0.79014291,  0.60754268,  0.63095296,
           0.41089627,  0.66372737,  0.99147137,  1.03829194,  1.34730772,
           1.38944624,  1.29580509,  1.0757484 ,  1.01956372]])
```

```
[ ]: test.shape
```

```
[ ]: (87, 24)
```

```
[ ]: test[:2]
```

```
[ ]: array([[ 1.0236215 ,  0.4564683 ,  0.43274588,  0.54917303,  0.52014428,
           0.53824823,  0.60847909,  0.52825984,  0.44179785,  0.32537069,
           0.27043456,  0.38654958,  0.59536933,  0.77235109,  0.89096321,
           0.95994552,  1.15908902,  1.29642937,  1.46498343,  1.47715677,
           1.54613909,  1.49744569,  1.30329638,  1.18093862],
          [ 0.12529078,  0.05037787, -0.04326328,  0.01760347, -0.09008385,
```

```
-0.17436088, -0.23522762, -0.19308911, -0.27268408, -0.38505346,
-0.40378168, -0.38973551, -0.38505346, -0.20245322, -0.05262739,
0.00355729, 0.07378815, 0.1674293 , 0.20956781, 0.31257307,
0.34534747, 0.3968501 , 0.24234221, 0.05974198]])
```

## 6 Create Input and Target values

The first 23 hours will be input to predict the 24th hour reading (i.e., target)

```
[ ]: # The first 23 columns (from 0 to 22) are inputs
```

```
train_inputs = train[:, :23]
```

```
pd.DataFrame(train_inputs, columns=np.arange(0,23,1))
```

```
[ ]:
```

	0	1	2	3	4	5	6	\
0	-1.986317	-1.836491	-1.850537	-1.756896	-1.906722	-1.948861	-1.925450	
1	0.401532	0.448353	0.485809	0.476445	0.499855	0.518584	0.541994	
2	-1.195049	-1.199731	-1.265280	-1.251234	-1.255916	-1.293373	-1.293373	
3	-1.616435	-1.602388	-1.635163	-1.653891	-1.761578	-1.799035	-1.799035	
4	0.982107	0.921241	0.696502	0.499855	0.401532	0.307891	0.270435	
..	...	...	...	...	...	...	...	
196	0.017603	-0.179043	-0.207135	-0.193089	-0.263320	-0.314823	-0.417828	
197	1.234938	1.099159	0.935287	0.799507	0.752686	0.748004	0.724594	
198	0.209568	0.312573	0.335983	0.335983	0.321937	0.247024	0.228296	
199	-0.001125	0.012921	-0.052627	-0.282048	-0.394418	-0.483377	-0.558290	
200	0.743322	0.663727	0.593497	0.462399	0.415578	0.424942	0.453035	

	7	8	9	...	13	14	15	\
0	-1.855219	-1.794353	-1.794353	...	-1.873948	-1.841173	-1.827127	
1	0.527948	0.541994	0.546676	...	0.607543	0.630953	0.410896	
2	-1.340193	-1.382332	-1.391696	...	-1.424470	-1.330829	-1.363603	
3	-1.799035	-1.850537	-1.883312	...	-1.948861	-1.911404	-1.958225	
4	0.247024	0.232978	0.209568	...	0.598179	0.762051	0.949333	
..	...	...	...	...	...	...	...	
196	-0.469330	-0.502105	-0.539561	...	-0.605110	-0.464648	-0.239910	
197	0.724594	0.701184	0.701184	...	1.089795	1.281759	1.436267	
198	0.214250	0.232978	0.275117	...	0.719912	0.949333	1.127251	
199	-0.633202	-0.684705	-0.722162	...	-0.427192	-0.164997	0.289163	
200	0.443671	0.382804	0.321937	...	0.640317	0.588814	0.701184	

	16	17	18	19	20	21	22
0	-1.808399	-1.742850	-1.658573	-1.555568	-1.490019	-1.461927	-1.452563
1	0.663727	0.991471	1.038292	1.347308	1.389446	1.295805	1.075748
2	-1.419788	-1.424470	-1.316783	-1.326147	-1.059270	-1.031177	-1.293373
3	-1.770942	-1.387014	-1.054588	-0.829849	-0.956264	-0.829849	-1.457245

```

4      1.057020  1.234938  1.295805  1.239620  1.202164  1.155343  1.206846
..      ...      ...      ...      ...      ...      ...
196    0.055060  0.396850  0.626271  0.776097  0.902512  0.804189  0.443671
197    1.623549  1.717190  1.867016  2.040252  2.124529  2.152622  2.040252
198    1.375400  1.637595  1.628231  1.562682  1.431585  1.333262  1.230256
199    0.527948  0.635635  0.930605  0.986789  1.024246  1.103841  0.888466
200    0.822917  0.836963  0.836963  0.949333  0.963379  1.089795  0.860374

```

[201 rows x 23 columns]

## 6.1 Add one more dimension to make it ready for RNNs

See here for more details: <https://keras.io/layers/recurrent/>, and <https://shivaverma.medium.com/understanding-input-and-output-shape-in-lstm-keras-c501ee95c65e>

```
[ ]: train_inputs
```

```
[ ]: array([[ -1.98631703e+00, -1.83649120e+00, -1.85053737e+00, ...,
           -1.49001896e+00, -1.46192662e+00, -1.45256251e+00],
          [ 4.01532160e-01,  4.48352732e-01,  4.85809190e-01, ...,
           1.38944624e+00,  1.29580509e+00,  1.07574840e+00],
          [-1.19504936e+00, -1.19973142e+00, -1.26528022e+00, ...,
           -1.05926970e+00, -1.03117735e+00, -1.29337256e+00],
          ...,
          [ 2.09567813e-01,  3.12573072e-01,  3.35983359e-01, ...,
           1.43158475e+00,  1.33326155e+00,  1.23025629e+00],
          [-1.12476250e-03,  1.29214092e-02, -5.26273921e-02, ...,
           1.02424577e+00,  1.10384075e+00,  8.88466112e-01],
          [ 7.43322338e-01,  6.63727365e-01,  5.93496507e-01, ...,
           9.63379028e-01,  1.08979457e+00,  8.60373769e-01]])
```

```
[ ]: train_inputs.shape
```

```
[ ]: (201, 23)
```

```
[ ]: #Create an additional dimension for train
```

```

train_x = train_inputs[:, :, np.newaxis]

train_x.shape

```

```
[ ]: (201, 23, 1)
```

```
[ ]: train_x
```

```
[ ]: array([[[ -1.98631703e+00],
             [-1.83649120e+00],
             [-1.85053737e+00],
```

```

...,
[-1.49001896e+00],
[-1.46192662e+00],
[-1.45256251e+00]],

[[ 4.01532160e-01],
 [ 4.48352732e-01],
 [ 4.85809190e-01],
...,
 [ 1.38944624e+00],
 [ 1.29580509e+00],
 [ 1.07574840e+00]],

[[-1.19504936e+00],
 [-1.19973142e+00],
 [-1.26528022e+00],
...,
 [-1.05926970e+00],
 [-1.03117735e+00],
 [-1.29337256e+00]],

...,

[[ 2.09567813e-01],
 [ 3.12573072e-01],
 [ 3.35983359e-01],
...,
 [ 1.43158475e+00],
 [ 1.33326155e+00],
 [ 1.23025629e+00]],

[[-1.12476250e-03],
 [ 1.29214092e-02],
 [-5.26273921e-02],
...,
 [ 1.02424577e+00],
 [ 1.10384075e+00],
 [ 8.88466112e-01]],

[[ 7.43322338e-01],
 [ 6.63727365e-01],
 [ 5.93496507e-01],
...,
 [ 9.63379028e-01],
 [ 1.08979457e+00],
 [ 8.60373769e-01]]])

```



## 6.2 Set the target

```
[ ]: # The last column (23) is TARGET

train_target = train[:, -1]

pd.DataFrame(train_target, columns=['23'])
```

```
[ ]:      23
0   -1.424470
1    1.019564
2   -1.485337
3   -1.962907
4    1.047656
..      ...
196  0.181475
197  0.410896
198  0.982107
199  0.560722
200  0.719912

[201 rows x 1 columns]
```

## 6.3 Repeat for TEST

```
[ ]: test.shape
```

```
[ ]: (87, 24)
```

```
[ ]: # The first 23 columns (from 0 to 22) are inputs

test_inputs = test[:, :23]
```

```
[ ]: # Create an additional dimension for test

test_x = test_inputs[:, :, np.newaxis]

test_x.shape
```

```
[ ]: (87, 23, 1)
```

```
[ ]: # The last column (23) is TARGET

test_target = test[:, -1]

pd.DataFrame(test_target, columns=['23'])
```

```
[ ]:          23
0    1.180939
1    0.059742
2   -1.003085
3    0.624086
4   -0.347597
..     ...
82  -1.892676
83    1.145979
84   -0.033899
85    1.057020
86    0.390607

[87 rows x 1 columns]
```

## 7 A normal (cross-sectional) NN

This model assumes that the data is NOT a time-series data set. It treats the data as cross-sectional and the columns being independent of each other.

```
[ ]: model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[23, 1]),
    keras.layers.Dense(23, activation='relu'),
    keras.layers.Dense(1, activation=None)

])

model.compile(loss="mse", optimizer='Adam')

history = model.fit(train_x, train_target, epochs=epoch_num)
```

```
/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
2024-04-09 07:30:43.848405: I metal_plugin/src/device/metal_device.cc:1154]
Metal device set to: Apple M1 Pro
2024-04-09 07:30:43.848428: I metal_plugin/src/device/metal_device.cc:296]
systemMemory: 16.00 GB
2024-04-09 07:30:43.848434: I metal_plugin/src/device/metal_device.cc:313]
maxCacheSize: 5.33 GB
2024-04-09 07:30:43.848452: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2024-04-09 07:30:43.848465: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
```

```
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)
```

Epoch 1/5

2024-04-09 07:30:44.428972: I

tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:117]

Plugin optimizer for device\_type GPU is enabled.

7/7 1s 24ms/step - loss:

0.7465

Epoch 2/5

7/7 0s 10ms/step - loss:

0.3339

Epoch 3/5

7/7 0s 8ms/step - loss:

0.2328

Epoch 4/5

7/7 0s 12ms/step - loss:

0.1573

Epoch 5/5

7/7 0s 9ms/step - loss:

0.1179

### 7.0.1 Predictions

```
[ ]: #Predict:
y_pred = model.predict(test_x)
```

3/3 0s 21ms/step

```
[ ]: # Remember, these are standardized values.

comparison = pd.DataFrame()

comparison['actual'] = scaler.inverse_transform([test_target]).flatten()
comparison['predicted'] = scaler.inverse_transform(y_pred).flatten()
```

```
[ ]: comparison
```

```
[ ]:      actual  predicted
0    86.021   83.915245
1    75.245   75.834129
2    65.030   71.919403
3    80.669   83.106102
4    71.330   78.152061
..     ...         ...
82    56.480   76.081505
83    85.685   85.226166
```

```
84 74.345 76.045578
85 84.830 85.395248
86 78.425 81.538994
```

```
[87 rows x 2 columns]
```


```
[ ]: mean_squared_error(comparison['actual'], comparison['predicted'])
```

```
[ ]: 22.148414581560083
```

```
[ ]: plt.plot(comparison['actual'], label = 'actual')
plt.plot(comparison['predicted'], label = 'predicted')

plt.legend()

plt.show()
```



## 8 Simple RNN

Simplest recurrent neural network

```
[ ]: model = keras.models.Sequential([
    keras.layers.SimpleRNN(32, activation='relu', input_shape=[23, 1]),
    keras.layers.Dense(1, activation=None)
])
```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
super().\_\_init\_\_(\*\*kwargs)

```
[ ]: np.random.seed(42)
    tf.random.set_seed(42)

    model.compile(loss="mse", optimizer='Adam')

    history = model.fit(train_x, train_target, epochs=epoch_num)
```

```
Epoch 1/5
7/7          3s 271ms/step - loss:
1.1885
Epoch 2/5
7/7          2s 250ms/step - loss:
0.9876
Epoch 3/5
7/7          2s 267ms/step - loss:
0.8507
Epoch 4/5
7/7          2s 255ms/step - loss:
0.7479
Epoch 5/5
7/7          2s 252ms/step - loss:
0.6585
```

### 8.0.1 Predictions

```
[ ]: #Predict:
    y_pred = model.predict(test_x)
```

```
3/3          0s 119ms/step
```

```
[ ]: #Remember, these are standardized values.

    comparison = pd.DataFrame()

    comparison['actual'] = scaler.inverse_transform([test_target]).flatten()
    comparison['predicted'] = scaler.inverse_transform(y_pred).flatten()
```

```
[ ]: comparison
```

```
[ ]:      actual  predicted
      0    86.021   80.769104
      1    75.245   76.220955
      2    65.030   74.551300
      3    80.669   80.229492
      4    71.330   75.179207
      ..      ...      ...
     82    56.480   76.151115
     83    85.685   80.316132
     84    74.345   76.672333
     85    84.830   80.800354
     86    78.425   76.670654
```

```
[87 rows x 2 columns]
```


```
[ ]: mean_squared_error(comparison['actual'], comparison['predicted'])
```

```
[ ]: 49.83867603891561
```

```
[ ]: plt.plot(comparison['actual'], label = 'actual')
      plt.plot(comparison['predicted'], label = 'predicted')

      plt.legend()

      plt.show()
```



## 8.1 Simple RNN with more layers

Be careful: when stacking RNN layers, you have to set “return\_sequences” to True. This enables the layer to send a “sequence” of values to the next layer – just like how it uses a sequence of values for training. However, if the output of RNN is sent to a DENSE layer, then a single value should be sent. That’s why there is no “return\_sequences” right before DENSE layers.

```
[ ]: model = keras.models.Sequential([
    keras.layers.SimpleRNN(32, activation='relu', return_sequences=True,
    ↪input_shape=[23, 1]),
    keras.layers.SimpleRNN(32, activation='relu', return_sequences=False),
    keras.layers.Dense(1, activation=None)
])
```

```
/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
```

```
[ ]: np.random.seed(42)
      tf.random.set_seed(42)

      model.compile(loss="mse", optimizer='Adam')

      history = model.fit(train_x, train_target, epochs=epoch_num)
```

```
Epoch 1/5
7/7          7s 675ms/step - loss:
0.8615
Epoch 2/5
7/7          5s 692ms/step - loss:
0.3967
Epoch 3/5
7/7          4s 621ms/step - loss:
0.1550
Epoch 4/5
7/7          5s 695ms/step - loss:
0.0933
Epoch 5/5
7/7          5s 669ms/step - loss:
0.0709
```

### 8.1.1 Predictions

```
[ ]: #Predict:
      y_pred = model.predict(test_x)
```

```
WARNING:tensorflow:5 out of the last 7 calls to <function
TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at
0x354e9f4c0> triggered tf.function retracing. Tracing is expensive and the
excessive number of tracings could be due to (1) creating @tf.function
repeatedly in a loop, (2) passing tensors with different shapes, (3) passing
Python objects instead of tensors. For (1), please define your @tf.function
outside of the loop. For (2), @tf.function has reduce_retracing=True option that
can avoid unnecessary retracing. For (3), please refer to
https://www.tensorflow.org/guide/function#controlling\_retracing and
https://www.tensorflow.org/api\_docs/python/tf/function for more details.
```

```
2/3          0s
132ms/stepWARNING:tensorflow:6 out of the last 9 calls to <function
TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at
0x354e9f4c0> triggered tf.function retracing. Tracing is expensive and the
excessive number of tracings could be due to (1) creating @tf.function
repeatedly in a loop, (2) passing tensors with different shapes, (3) passing
Python objects instead of tensors. For (1), please define your @tf.function
outside of the loop. For (2), @tf.function has reduce_retracing=True option that
can avoid unnecessary retracing. For (3), please refer to
https://www.tensorflow.org/guide/function#controlling\_retracing and
```



[https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) for more details.  
3/3 1s 225ms/step

```
[ ]: #Remember, these are standardized values.
```

```
comparison = pd.DataFrame()

comparison['actual'] = scaler.inverse_transform([test_target]).flatten()
comparison['predicted'] = scaler.inverse_transform(y_pred).flatten()
```

```
[ ]: comparison
```

```
[ ]:      actual  predicted
0    86.021   83.467468
1    75.245   75.912086
2    65.030   70.674118
3    80.669   82.433907
4    71.330   74.155128
..     ...      ...
82   56.480   59.918369
83   85.685   83.932892
84   74.345   75.744240
85   84.830   83.924927
86   78.425   77.446617
```

```
[87 rows x 2 columns]
```


```
[ ]: mean_squared_error(comparison['actual'], comparison['predicted'])
```

```
[ ]: 5.183264379163942
```

```
[ ]: plt.plot(comparison['actual'], label = 'actual')
plt.plot(comparison['predicted'], label = 'predicted')

plt.legend()

plt.show()
```



## 9 LSTM with one layer

```
[ ]: model = keras.models.Sequential([
    keras.layers.LSTM(32, activation='relu', input_shape=[23, 1]),
    keras.layers.Dense(1, activation=None)
])
```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
super().\_\_init\_\_(\*\*kwargs)

```
[ ]: np.random.seed(42)
    tf.random.set_seed(42)

    model.compile(loss="mse", optimizer='Adam')

    history = model.fit(train_x, train_target, epochs=epoch_num)
```

Epoch 1/5

7/7

5s 440ms/step - loss:

```
1.2531
Epoch 2/5
7/7          3s 487ms/step - loss:
1.0240
Epoch 3/5
7/7          3s 479ms/step - loss:
0.8281
Epoch 4/5
7/7          3s 481ms/step - loss:
0.6236
Epoch 5/5
7/7          3s 467ms/step - loss:
0.5775
```

### 9.0.1 Predictions

```
[ ]: #Predict:
y_pred = model.predict(test_x)
```

```
3/3          1s 154ms/step
```

```
[ ]: #Remember, these are standardized values.

comparison = pd.DataFrame()

comparison['actual'] = scaler.inverse_transform([test_target]).flatten()
comparison['predicted'] = scaler.inverse_transform(y_pred).flatten()
```


```
[ ]: mean_squared_error(comparison['actual'], comparison['predicted'])
```

```
[ ]: 77.44127969120099
```

```
[ ]: plt.plot(comparison['actual'], label = 'actual')
plt.plot(comparison['predicted'], label = 'predicted')

plt.legend()

plt.show()
```



## 10 LSTM with more layers

```
[ ]: model = keras.models.Sequential([
    keras.layers.LSTM(32, activation='tanh', return_sequences=True,
    ↪input_shape=[23, 1]),
    keras.layers.LSTM(32, activation='tanh', return_sequences=False),
    keras.layers.Dense(1, activation=None)
])
```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
 super().\_\_init\_\_(\*\*kwargs)

```
[ ]: np.random.seed(42)
    tf.random.set_seed(42)

    model.compile(loss="mse", optimizer='Adam')

    history = model.fit(train_x, train_target, epochs=epoch_num)
```

```
Epoch 1/5
7/7          2s 31ms/step - loss:
0.7669
Epoch 2/5
7/7          0s 16ms/step - loss:
0.2524
Epoch 3/5
7/7          0s 14ms/step - loss:
0.1259
Epoch 4/5
7/7          0s 15ms/step - loss:
0.1136
Epoch 5/5
7/7          0s 14ms/step - loss:
0.0880
```

### 10.0.1 Predictions

```
[ ]: #Predict:
y_pred = model.predict(test_x)
```

```
3/3          0s 101ms/step
```

```
[ ]: #Remember, these are standardized values.

comparison = pd.DataFrame()

comparison['actual'] = scaler.inverse_transform([test_target]).flatten()
comparison['predicted'] = scaler.inverse_transform(y_pred).flatten()
```


```
[ ]: mean_squared_error(comparison['actual'], comparison['predicted'])
```

```
[ ]: 12.661517235157207
```

```
[ ]: plt.plot(comparison['actual'], label = 'actual')
plt.plot(comparison['predicted'], label = 'predicted')

plt.legend()

plt.show()
```



## 11 GRU (with more layers)

```
[ ]: model = keras.models.Sequential([
    keras.layers.GRU(32, activation='relu', return_sequences=True,
    ↪input_shape=[23, 1]),
    keras.layers.GRU(32, activation='relu', return_sequences=False),
    keras.layers.Dense(1, activation=None)
])
```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
super().\_\_init\_\_(\*\*kwargs)

```
[ ]: np.random.seed(42)
    tf.random.set_seed(42)

    model.compile(loss="mse", optimizer='RMSprop')

    history = model.fit(train_x, train_target, epochs=epoch_num)
```

```
Epoch 1/5
7/7          13s 2s/step - loss:
0.8036
Epoch 2/5
7/7          11s 2s/step - loss:
0.5600
Epoch 3/5
7/7          11s 2s/step - loss:
0.4047
Epoch 4/5
7/7          11s 2s/step - loss:
0.2556
Epoch 5/5
7/7          11s 2s/step - loss:
0.1227
```

### 11.0.1 Predictions

```
[ ]: #Predict:
y_pred = model.predict(test_x)
```

```
3/3          1s 358ms/step
```

```
[ ]: #Remember, these are standardized values.

comparison = pd.DataFrame()

comparison['actual'] = scaler.inverse_transform([test_target]).flatten()
comparison['predicted'] = scaler.inverse_transform(y_pred).flatten()
```


```
[ ]: mean_squared_error(comparison['actual'], comparison['predicted'])
```

```
[ ]: 5.907716083886514
```

```
[ ]: plt.plot(comparison['actual'], label = 'actual')
plt.plot(comparison['predicted'], label = 'predicted')

plt.legend()

plt.show()
```



## 12 Conv1D

### 12.0.1 Last Layer: GRU (you can change it to SimpleRNN or LSTM as well)

```
[ ]: model = keras.models.Sequential([
    keras.layers.Conv1D(filters=20, kernel_size=3, strides=1, padding="valid",
↪input_shape=[23, 1]),
    keras.layers.GRU(32, activation='relu', return_sequences=True),
    keras.layers.GRU(32, activation='relu', return_sequences=False),
    keras.layers.Dense(1, activation=None)
])
```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-packages/keras/src/layers/convolutional/base\_conv.py:99: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
    super().__init__(
```

```
[ ]: np.random.seed(42)
      tf.random.set_seed(42)
```



```
model.compile(loss="mse", optimizer='Adam')

history = model.fit(train_x, train_target, epochs=epoch_num)
```

```
Epoch 1/5
7/7          14s 1s/step - loss:
0.9220
Epoch 2/5
7/7          10s 1s/step - loss:
0.7516
Epoch 3/5
7/7          10s 1s/step - loss:
0.6433
Epoch 4/5
7/7          10s 1s/step - loss:
0.5135
Epoch 5/5
7/7          10s 1s/step - loss:
0.3517
```

## 12.0.2 Predictions

```
[ ]: #Predict:
y_pred = model.predict(test_x)
```

```
3/3          1s 325ms/step
```

```
[ ]: #Remember, these are standardized values.

comparison = pd.DataFrame()

comparison['actual'] = scaler.inverse_transform([test_target]).flatten()
comparison['predicted'] = scaler.inverse_transform(y_pred).flatten()
```

```
[ ]: mean_squared_error(comparison['actual'], comparison['predicted'])
```

```
[ ]: 18.93950135820779
```

## 13 Forecasting Several Steps Ahead

### 13.1 Now let's create an RNN that predicts 12 next values at once:

```
[ ]: # The first 12 columns (from 0 to 11) are inputs

train_inputs = train[:, :12]

pd.DataFrame(train_inputs, columns=np.arange(0,12,1))
```

```
[ ]:
```

	0	1	2	3	4	5	6	\
0	-1.986317	-1.836491	-1.850537	-1.756896	-1.906722	-1.948861	-1.925450	
1	0.401532	0.448353	0.485809	0.476445	0.499855	0.518584	0.541994	
2	-1.195049	-1.199731	-1.265280	-1.251234	-1.255916	-1.293373	-1.293373	
3	-1.616435	-1.602388	-1.635163	-1.653891	-1.761578	-1.799035	-1.799035	
4	0.982107	0.921241	0.696502	0.499855	0.401532	0.307891	0.270435	
..	...	...	...	...	...	...	...	
196	0.017603	-0.179043	-0.207135	-0.193089	-0.263320	-0.314823	-0.417828	
197	1.234938	1.099159	0.935287	0.799507	0.752686	0.748004	0.724594	
198	0.209568	0.312573	0.335983	0.335983	0.321937	0.247024	0.228296	
199	-0.001125	0.012921	-0.052627	-0.282048	-0.394418	-0.483377	-0.558290	
200	0.743322	0.663727	0.593497	0.462399	0.415578	0.424942	0.453035	

	7	8	9	10	11
0	-1.855219	-1.794353	-1.794353	-1.761578	-1.770942
1	0.527948	0.541994	0.546676	0.518584	0.574768
2	-1.340193	-1.382332	-1.391696	-1.401060	-1.424470
3	-1.799035	-1.850537	-1.883312	-1.939496	-1.972271
4	0.247024	0.232978	0.209568	0.200204	0.167429
..	...	...	...	...	...
196	-0.469330	-0.502105	-0.539561	-0.548925	-0.516151
197	0.724594	0.701184	0.701184	0.691820	0.738640
198	0.214250	0.232978	0.275117	0.284481	0.211128
199	-0.633202	-0.684705	-0.722162	-0.754936	-0.797074
200	0.443671	0.382804	0.321937	0.331301	0.429625

[201 rows x 12 columns]

```
[ ]: #Create an additional dimension for train
```

```
train_x = train_inputs.reshape(201,12,1)

train_x.shape
```

```
[ ]: (201, 12, 1)
```

```
[ ]: # The last 12 readings (from 12 to 23) are TARGET
```

```
train_target = train[:, -12:]

pd.DataFrame(train_target, columns=np.arange(12,24,1))
```

```
[ ]:
```

	12	13	14	15	16	17	18	\
0	-1.817763	-1.873948	-1.841173	-1.827127	-1.808399	-1.742850	-1.658573	
1	0.790143	0.607543	0.630953	0.410896	0.663727	0.991471	1.038292	
2	-1.415106	-1.424470	-1.330829	-1.363603	-1.419788	-1.424470	-1.316783	
3	-2.000363	-1.948861	-1.911404	-1.958225	-1.770942	-1.387014	-1.054588	

4	0.415578	0.598179	0.762051	0.949333	1.057020	1.234938	1.295805
..	...	...	...	...	...	...	...
196	-0.614474	-0.605110	-0.464648	-0.239910	0.055060	0.396850	0.626271
197	0.935287	1.089795	1.281759	1.436267	1.623549	1.717190	1.867016
198	0.401532	0.719912	0.949333	1.127251	1.375400	1.637595	1.628231
199	-0.670659	-0.427192	-0.164997	0.289163	0.527948	0.635635	0.930605
200	0.654363	0.640317	0.588814	0.701184	0.822917	0.836963	0.836963

	19	20	21	22	23
0	-1.555568	-1.490019	-1.461927	-1.452563	-1.424470
1	1.347308	1.389446	1.295805	1.075748	1.019564
2	-1.326147	-1.059270	-1.031177	-1.293373	-1.485337
3	-0.829849	-0.956264	-0.829849	-1.457245	-1.962907
4	1.239620	1.202164	1.155343	1.206846	1.047656
..	...	...	...	...	...
196	0.776097	0.902512	0.804189	0.443671	0.181475
197	2.040252	2.124529	2.152622	2.040252	0.410896
198	1.562682	1.431585	1.333262	1.230256	0.982107
199	0.986789	1.024246	1.103841	0.888466	0.560722
200	0.949333	0.963379	1.089795	0.860374	0.719912

[201 rows x 12 columns]

### 13.2 Repeat for TEST

```
[ ]: # The first 12 columns (from 0 to 11) are inputs

test_inputs = test[:, :12]

pd.DataFrame(test_inputs, columns=np.arange(0,12,1))
```

[ ]:	0	1	2	3	4	5	6	\
0	1.023621	0.456468	0.432746	0.549173	0.520144	0.538248	0.608479	
1	0.125291	0.050378	-0.043263	0.017603	-0.090084	-0.174361	-0.235228	
2	-0.445920	-0.497423	-0.516151	-0.511469	-0.544243	-0.614474	-0.637885	
3	0.791704	0.651554	0.591936	0.510780	0.529196	0.490803	0.448353	
4	-0.403782	-0.417828	-0.455284	-0.544243	-0.619156	-0.656613	-0.698751	
..	...	...	...	...	...	...	...	
82	-0.047945	-0.076038	-0.108812	-0.155633	-0.174361	-0.146269	-0.090084	
83	0.982107	0.907194	0.925923	0.701184	0.574768	0.513902	0.438989	
84	-0.202453	-0.272684	-0.347597	-0.342915	-0.445920	-0.534879	-0.600428	
85	1.038292	0.982107	0.958697	0.916558	0.846328	0.841646	0.804189	
86	0.768293	0.689947	0.624398	0.571959	0.526387	0.453971	0.464584	
	7	8	9	10	11			
0	0.528260	0.441798	0.325371	0.270435	0.386550			
1	-0.193089	-0.272684	-0.385053	-0.403782	-0.389736			

```

2  -0.609792 -0.619156 -0.651931 -0.605110 -0.708115
3   0.331613  0.312885  0.323810  0.280111  0.374688
4  -0.759618 -0.815803 -0.867305 -0.909444 -0.923490
..      ...      ...      ...      ...      ...
82 -0.076038 -0.450602 -0.750254 -0.839213 -0.974993
83  0.406214  0.387486  0.387486  0.373440  0.392168
84 -0.633202 -0.628520 -0.670659 -0.689387 -0.708115
85  0.752686  0.752686  0.724594  0.663727  0.743322
86  0.412457  0.359394  0.273244  0.152135  0.148389

```

[87 rows x 12 columns]

```
[ ]: #Create an additional dimension for test
```

```

test_x = test_inputs.reshape(87,12,1)

test_x.shape

```

```
[ ]: (87, 12, 1)
```

```
[ ]: # The last 12 columns are TARGET
```

```

test_target = test[:, -12:]

pd.DataFrame(test_target, columns=np.arange(12,24,1))

```

```

[ ]:
      12      13      14      15      16      17      18  \
0  0.595369  0.772351  0.890963  0.959946  1.159089  1.296429  1.464983
1 -0.385053 -0.202453 -0.052627  0.003557  0.073788  0.167429  0.209568
2 -0.825167 -0.806439 -0.431874 -0.352279 -0.427192 -0.310141 -0.188407
3  0.583820  0.839461  1.129124  1.293932  1.375400  1.511804  1.636659
4 -0.909444 -1.007767 -0.960946 -0.698751 -0.314823  0.415578  0.560722
..      ...      ...      ...      ...      ...      ...      ...
82 -1.195049 -1.312101 -1.227824 -1.049906 -1.092044 -1.096726 -1.007767
83  0.626271  0.715230  0.916558  1.136615  1.469041  1.651641  1.796785
84 -0.764300 -0.661295 -0.352279  0.026968  0.331301  0.509219  0.654363
85  0.265752  0.893148  1.267713  1.108523  1.314533  1.628231  1.548636
86  0.182412  0.219556  0.229857  0.283232  0.426503  0.458029  0.474884

      19      20      21      22      23
0  1.477157  1.546139  1.497446  1.303296  1.180939
1  0.312573  0.345347  0.396850  0.242342  0.059742
2  0.031650 -0.033899 -0.108812 -0.436556 -1.003085
3  1.641965  1.550862  1.386013  1.051090  0.624086
4  0.659045  0.340665  0.069106 -0.169679 -0.347597
..      ...      ...      ...      ...      ...
82 -1.026495 -1.106090 -1.241870 -1.630481 -1.892676

```

```

83  1.515862  1.450313  1.220892  1.384764  1.145979
84  0.630953  0.598179  0.457717  0.223614 -0.033899
85  1.600139  1.454995  1.511180  1.277077  1.057020
86  0.409960  0.395289  0.421509  0.435555  0.390607

```

[87 rows x 12 columns]

## 14 GRU

```

[ ]: model = keras.models.Sequential([
    keras.layers.GRU(32, activation='relu', return_sequences=True,
    ↪input_shape=[12, 1]),
    keras.layers.GRU(32, activation='relu', return_sequences=False),
    keras.layers.Dense(12, activation=None)
])

```

```

/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

```

```

[ ]: np.random.seed(42)
    tf.random.set_seed(42)

    model.compile(loss="mse", optimizer='Adam')

    history = model.fit(train_x, train_target, epochs=epoch_num)

```

```

Epoch 1/5
7/7          7s 536ms/step - loss:
1.1441
Epoch 2/5
7/7          3s 474ms/step - loss:
1.0978
Epoch 3/5
7/7          3s 487ms/step - loss:
1.0587
Epoch 4/5
7/7          3s 488ms/step - loss:
1.0147
Epoch 5/5
7/7          3s 459ms/step - loss:
0.9574

```

### 14.0.1 Predictions

```
[ ]: #Predict:
y_pred = model.predict(test_x)
```

3/3 1s 232ms/step

```
[ ]: #Remember, these are standardized values.

actual = pd.DataFrame(scaler.inverse_transform(test_target))
predicted = pd.DataFrame(scaler.inverse_transform(y_pred))
```

```
[ ]: actual
```

```
[ ]:
      0      1      2      3      4      5      6      7      8  \
0  80.393  82.094  83.234  83.897  85.811  87.131  88.751  88.868  89.531000
1  70.970  72.725  74.165  74.705  75.380  76.280  76.685  77.675  77.990000
2  66.740  66.920  70.520  71.285  70.565  71.690  72.860  74.975  74.345000
3  80.282  82.739  85.523  87.107  87.890  89.201  90.401  90.452  89.576393
4  65.930  64.985  65.435  67.955  71.645  78.665  80.060  81.005  77.945000
..    ...    ...    ...    ...    ...    ...    ...    ...
82  63.185  62.060  62.870  64.580  64.175  64.130  64.985  64.805  64.040000
83  80.690  81.545  83.480  85.595  88.790  90.545  91.940  89.240  88.610000
84  67.325  68.315  71.285  74.930  77.855  79.565  80.960  80.735  80.420000
85  77.225  83.255  86.855  85.325  87.305  90.320  89.555  90.050  88.655000
86  76.424  76.781  76.880  77.393  78.770  79.073  79.235  78.611  78.470000

      9      10      11
0  89.063  87.197  86.021
1  78.485  77.000  75.245
2  73.625  70.475  65.030
3  87.992  84.773  80.669
4  75.335  73.040  71.330
..    ...    ...    ...
82  62.735  59.000  56.480
83  86.405  87.980  85.685
84  79.070  76.820  74.345
85  89.195  86.945  84.830
86  78.722  78.857  78.425
```

[87 rows x 12 columns]

```
[ ]: predicted
```

```
[ ]:
      0      1      2      3      4      5  \
0  75.710274  76.159607  75.889503  76.173485  74.491684  74.442482
1  73.025398  73.948112  73.579193  75.142609  74.772263  75.695740
2  72.061836  73.643280  72.786995  74.916153  74.808380  75.914146
```

3	75.603302	76.022293	75.804817	76.094887	74.528267	74.488770
4	71.479134	73.488319	72.309883	74.831909	74.780487	76.029640
..	...	...	...	...	...	...
82	72.156029	73.719513	72.835388	74.997971	74.731590	75.878960
83	75.752594	76.221657	75.922379	76.189705	74.491486	74.412437
84	72.015129	73.631264	72.741776	74.907677	74.806061	75.923180
85	76.355865	76.999901	76.385094	76.680244	74.223328	74.172867
86	75.486946	75.841896	75.727425	76.083000	74.508736	74.581299

	6	7	8	9	10	11
0	76.356964	76.361732	76.299683	77.409576	75.419708	77.005798
1	74.771301	74.884758	75.348892	75.261559	74.148224	75.608627
2	74.319107	74.430626	75.067566	74.989601	73.326004	75.468018
3	76.259109	76.263405	76.222855	77.234825	75.389908	76.862595
4	74.029396	74.143074	74.919838	74.866920	72.793411	75.389748
..	...	...	...	...	...	...
82	74.342766	74.434380	75.128151	75.062981	73.335808	75.460693
83	76.390038	76.391083	76.312370	77.477257	75.431549	77.045738
84	74.293282	74.398834	75.052490	74.976540	73.278160	75.452759
85	76.962563	76.986366	76.803223	78.439301	75.542992	77.886246
86	76.179092	76.197670	76.219551	77.085602	75.360512	76.803185

[87 rows x 12 columns]

```
[ ]: mean_squared_error(actual, predicted)
```


```
[ ]: 78.06923496947906
```

```
[ ]: # Plot a random row to see the accuracy of predictions

random_row = np.random.randint(low=0, high=86)

plt.plot(actual.iloc[random_row], label='actual')
plt.plot(predicted.iloc[random_row], label='predicted')

plt.legend()
plt.show()
```



## 15 Sliding window

Prior 18 hours predicts next 6 hours

```
[ ]: steps_for_prediction = 18
     steps_to_predict = 6

     #Be careful: sums to 24 hours
```

```
[ ]: train
```

```
[ ]: array([[ -1.98631703e+00,  -1.83649120e+00,  -1.85053737e+00, ...,
           -1.46192662e+00,  -1.45256251e+00,  -1.42447016e+00],
          [  4.01532160e-01,   4.48352732e-01,   4.85809190e-01, ...,
           1.29580509e+00,   1.07574840e+00,   1.01956372e+00],
          [-1.19504936e+00,  -1.19973142e+00,  -1.26528022e+00, ...,
           -1.03117735e+00,  -1.29337256e+00,  -1.48533691e+00],
          ...,
          [  2.09567813e-01,   3.12573072e-01,   3.35983359e-01, ...,
           1.33326155e+00,   1.23025629e+00,   9.82107257e-01],
          [-1.12476250e-03,   1.29214092e-02,  -5.26273921e-02, ...,
           1.10384075e+00,   8.88466112e-01,   5.60722106e-01],
```



```
[ 7.43322338e-01,  6.63727365e-01,  5.93496507e-01, ...,
 1.08979457e+00,  8.60373769e-01,  7.19912052e-01]])
```

```
[ ]: train.flatten().shape
```

```
[ ]: (4824,)
```

```
[ ]: train_inputs_sw = []
train_target_sw = []

for i in range(0,4824-24):
    input_row = train.flatten()[i:i+steps_for_prediction]
    target_row = train.flatten()[i+steps_for_prediction:
    ↪ i+steps_for_prediction+steps_to_predict]
    train_inputs_sw.append((input_row))
    train_target_sw.append((target_row))
```

```
[ ]: train_inputs = np.vstack(train_inputs_sw)

train_targets = np.vstack(train_target_sw)
```

```
[ ]: train_targets.shape
```

```
[ ]: (4800, 6)
```

```
[ ]: # Repeat for test

test_inputs_sw = []
test_target_sw = []

for i in range(0,test.flatten().shape[0]-24):
    input_row = test.flatten()[i:i+steps_for_prediction]
    target_row = test.flatten()[i+steps_for_prediction:
    ↪ i+steps_for_prediction+steps_to_predict]
    test_inputs_sw.append((input_row))
    test_target_sw.append((target_row))

test_inputs = np.vstack(test_inputs_sw)

test_targets = np.vstack(test_target_sw)
```

## 16 GRU

```
[ ]: model = keras.models.Sequential([
    keras.layers.GRU(32, activation='relu', return_sequences=True,
    ↪ input_shape=[18, 1]),
    keras.layers.GRU(32, activation='relu', return_sequences=False),
```

```
keras.layers.Dense(steps_to_predict, activation=None)
])
```

```
/Users/timsmith/miniconda3/envs/dsp/lib/python3.11/site-
packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(**kwargs)
```

```
[ ]: np.random.seed(42)
      tf.random.set_seed(42)

      model.compile(loss="mse", optimizer='Adam')

      history = model.fit(train_inputs, train_targets, epochs=epoch_num)
```

```
Epoch 1/5
150/150      85s 549ms/step -
loss: 0.7335
Epoch 2/5
150/150      83s 555ms/step -
loss: 0.2851
Epoch 3/5
150/150      77s 513ms/step -
loss: 0.2646
Epoch 4/5
150/150      54s 362ms/step -
loss: 0.2550
Epoch 5/5
150/150      45s 299ms/step -
loss: 0.2476
```

### 16.0.1 Predictions

```
[ ]: #Predict:
      y_pred = model.predict(test_inputs)
```

```
65/65      4s 57ms/step
```

```
[ ]: #Remember, these are standardized values.

      actual = pd.DataFrame(scaler.inverse_transform(test_targets))
      predicted = pd.DataFrame(scaler.inverse_transform(y_pred))
```

```
[ ]: actual
```

```
[ ]:
      0      1      2      3      4      5
0    88.751  88.868  89.531  89.063  87.197  86.021
1    88.868  89.531  89.063  87.197  86.021  75.875
```

```

2      89.531  89.063  87.197  86.021  75.875  75.155
3      89.063  87.197  86.021  75.875  75.155  74.255
4      87.197  86.021  75.875  75.155  74.255  74.840
...
2059   76.781  76.880  77.393  78.770  79.073  79.235
2060   76.880  77.393  78.770  79.073  79.235  78.611
2061   77.393  78.770  79.073  79.235  78.611  78.470
2062   78.770  79.073  79.235  78.611  78.470  78.722
2063   79.073  79.235  78.611  78.470  78.722  78.857

```

[2064 rows x 6 columns]

```
[ ]: predicted
```

```

[ ]:
      0      1      2      3      4      5
0      87.054977  87.487541  85.834480  85.526680  84.404007  83.329414
1      88.089355  87.415459  84.871231  84.101242  82.409752  80.780205
2      87.591972  85.875031  83.330467  81.639549  79.749985  78.018005
3      87.657585  85.078636  82.162445  80.146370  78.001503  76.134148
4      86.442322  83.334702  80.434105  77.983383  75.658669  74.115738
...
2059   76.865219  78.937889  80.954544  82.301857  84.409531  85.125687
2060   77.326805  79.606239  81.456772  82.817108  84.677170  85.337708
2061   77.405602  79.738190  81.370369  82.691170  84.216827  84.952332
2062   77.916618  80.251701  81.486099  82.808311  83.912796  84.648819
2063   79.387184  81.782616  82.611473  83.948517  84.647949  85.253090

```

[2064 rows x 6 columns]

```
[ ]: mean_squared_error(actual, predicted)
```

```
[ ]: 24.017596273764052
```

```


[ ]: # Plot a random row to see the accuracy of predictions

random_row = np.random.randint(low=0, high=2063)

plt.plot(actual.iloc[random_row], label='actual')
plt.plot(predicted.iloc[random_row], label='predicted')

plt.legend()
plt.show()

```



16.1 We could try using 6 steps to predict the next 6 steps (maybe 12 steps is too long)

[ ]: